

Computer Architecture. Week 10

Overview of Modern Computers Vulnerabilities

Alexander Tormasov

Innopolis University

a.tormasov@innopolis.ru

October 28, 2021

Content of the class

- Vulnerabilities and Exploits
- Security Goals
- Meltdown and Spectre
- Threat-o-meter
- Memory model and process
- Side channeling
- Execution order – out-of-order execution
- Speculative execution
- Summary

Vulnerabilities and Exploits

- **Vulnerability** is a weakness in a system.
- An **exploit** is an attack that leverages that vulnerability.
- So vulnerable means there is theoretically a way to exploit something (i.e., a **vulnerability exists**), exploitable means that there is a definite path to doing so in the wild.
- Naturally, attackers want to find weaknesses that are actually exploitable.


Vulnerable CPU

- The current set of **vulnerabilities exploit** how our modern processors work.
- Unlike software-based attacks, these **hardware vulnerabilities** allow programs to **steal data** which is currently processed on the computer.
- So, the problem is very low level and uses the actual **CPU hardware architecture**.


Security Goals

- A process must be isolated (Protected) from other processes!

Security Goals: App Isolation



You don't want
this




To read
that

Security Goals: Browser Tab Isolation




You don't want
this



To read
that

Security Goals: Cloud Isolation



You don't want
this

To read
that

Security Goals: Software vulnerabilities

- The most common software security vulnerabilities include:
 - Missing data encryption
 - OS command injection
 - SQL injection
 - Buffer overflow
 - Missing authentication for critical function
 - Missing authorization
 - Unrestricted upload of dangerous file types
 - Reliance on untrusted inputs in a security decision
 - Cross-site scripting and forgery
 - Download of codes without integrity checks
 - Use of broken algorithms
 - URL redirection to untrusted sites
 - Path traversal
 - Bugs
 - Weak passwords
 - Software that is already infected with virus

Security Goals: Memory Isolation (1/2)



PROCESS A



PROCESS B

You don't want
this

To read
that

Security Goals: Memory Isolation (2/2)

- The more a system relies on process isolation to achieve its security goal, the more critical **Meltdown and Spectre** are.
- Fortunately an attacker must be able to **execute his code on a system** to exploit the Meltdown and Spectre attacks.

Security Goals: Cold Boot Attacks

On February 2008...



Hackers Decrypt Computer by Freezing Memory Chips

TECHNOLOGY

Researchers Find Way to Steal Encrypted Data

Security Goals: Cold Boot Attacks

- Memory modules gradually lose data over time as they lose power, but do not immediately lose all data when power is lost
- Depending on temperature and environmental conditions, memory modules can potentially retain, at least, some data for up to 90 minutes after power loss.
- With certain memory modules, the time window for an attack can be extended to hours or even weeks by cooling them with freeze spray.
- Liquid nitrogen, freeze spray or compressed air cans can be improvised to cool memory modules, and thereby slow down the degradation of volatile memory

Security Goals: Cold Boot Attacks


Cold boot attacks can steal encryption keys from nearly any laptop

1.



Attacker gets physical access to a company laptop

2.




Attacker manipulates firmware settings

3.



Attacker performs cold reboot
Into USB key

4.



Attacker gets encryption keys
from memory


Security Goals: Cold Boot Attacks

- A cold boot attack is a type of **side channel attack** in which an attacker with **physical access** to a computer performs a memory dump of a **computer's random access memory** by performing a hard reset of the target machine.
- Typically, cold boot attacks are used to **retrieve encryption keys** from a running operating system for malicious or criminal investigative reasons.
- The attack relies on the data remanence property of DRAM and SRAM to **retrieve memory contents** that remain readable in the seconds to minutes after power has been removed

Foreshadow (Security vulnerability)

- Foreshadow is a vulnerability that affects modern microprocessors that was first discovered in [January 2018](#).
- The vulnerability is a [speculative execution attack](#) on Intel processors that may result in the disclosure of sensitive information stored in personal computers and third-party clouds.

Meltdown and Spectre

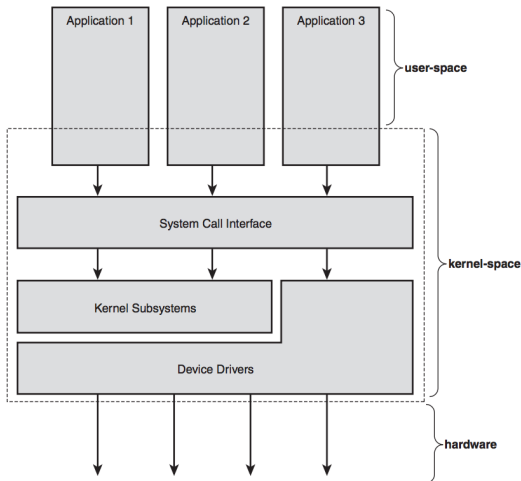


- Two security flaws were unveiled by security researchers (3rd January 2018)
 - Meltdown
 - Spectre
- At Google's Project Zero in conjunction with academic and industry researchers from several countries

Meltdown and Spectre

- The vulnerabilities behind the devastating Meltdown and Spectre attacks have existed for **decades**.
- Meltdown is a massive vulnerability on nearly every **Intel chip made since 1995**, but is largely being fixed with software patches.
- Spectre is more **difficult to exploit**, but will likely be with us for years
- Both exploits allow a **malicious user to access data**, whether that's your password, credit-card number, or just your personal photos stored on a cloud server.

Memory – User and Kernel Spaces



NOTE: More details in Operating System (OS) course.

The Attacker Code – Example

```

1 ; rcx = kernel address
2 ; rbx = probe array
3 retry:
4 mov al, byte [rcx]
5 shl rax, 0xc
6 jz retry
7 mov rbx, qword [rbx + rax]

```


- Explanation:** An **inaccessible kernel address** is moved to a register, **raising an exception**. The subsequent instructions are already executed **out of order** before the exception is raised, **leaking the content of the kernel address through the indirect memory access**.

Meltdown




- **Result:** Programs can read memory it should not
- **Affects:** All modern CPU/OS
- **Vector:** Uses out of order execution to read forbidden memory and cache timing as side channel to exfiltrate data
- **How bad:** Bad
- **Fixes:** Needs changes in CPU and/or OS patches. Modest (X%) to severe (XX%) performance impact, higher on older CPU. Performance impact varies and depends on CPU and workload type.

Spectre



- **Result:** Programs can read all memory
- **Affects:** All modern CPU/OS
- **Vector:** Uses speculative execution to read forbidden memory and cache timing to exfiltrate data
- **How bad:** Very bad
- **Fixes:** Needs changes in CPU and/or changes in programs. Performance impact varies and depends on CPU and workload type.

Threat-o-meter




Process

- A process is an **instance of a computer program** that is being executed.
- It contains the program **code** and its **activity**.
- Depending on the **Operating System** (OS), a process may be made up of multiple threads of execution that execute instructions concurrently.


Process – Components

- When a program is loaded into the memory and it becomes a process, it can be divided into **four sections** – stack, heap, text and data.




NOTE: More details in OS course.

Memory Model (1/5)




- CPU and OS **isolate processes memory** from each other
- Virtual Memory gives each process its **own address space**
- Each **address space** starts at “virtual address 0x000..0”

Memory Model (2/5)




- Memory is **split into pages** (each 4KiB on x86)
- The kernel maps its own **memory into each process**
- This “kernel” memory is **only accessible by the kernel**

Memory Model (3/5)



- b) and c) are **completely different error scenarios**
 - c) Kernel memory pages are marked “**kernel only**” but the process could try to access the pages via a **pointer**
 - b) Process **b** has **no possibility** to even describe the address


Memory Model (4/5)



- c) is vulnerable to Meltdown and Spectre
- b) is vulnerable to Spectre


Memory Model – Virtual Memory (1/3)

- Virtual memory is backed by physical RAM
- Virtual memory is much, much larger than physical RAM
- Not all virtual memory is backed by RAM



- Like a matryoshka doll the kernel maps all physical memory into its address space
- Reading kernel memory allows reading of all (mapped) memory of all processes


Memory Model – Virtual Memory (2/3)



The system works because principle of locality holds.

Thrashing : System swaps in/out all the time, no real work is done.

Memory Model – Virtual Memory (2/3)



In virtual memory, **blocks of memory (called pages)** are mapped from one set of addresses (called virtual addresses) to another set (called physical addresses)

Memory Model (5/5)

Virtual memory map with 4 level page tables:

Start addr	Offset	End addr	Size	VM area description
0000000000000000	0	00007fffffffff	128 TB	user-space virtual memory, different per mm
0000800000000000	+128 TB	ffff7fffffffff	~16M TB	... huge, almost 64 bits wide hole of non-canonical virtual memory addresses up to the -128 TB starting offset of kernel mappings.
				Kernel-space virtual memory, shared between all processes:
ffff800000000000	-128 TB	ffff87fffffffff	8 TB	... guard hole, also reserved for hypervisor
ffff880000000000	-120 TB	ffffc7fffffffff	64 TB	direct mapping of all physical memory (page_offset_base)
ffffc80000000000	-56 TB	ffffc87fffffffff	1 TB	... unused hole
ffffc90000000000	-55 TB	ffffc97fffffffff	32 TB	vmalloc/ioremap space (vmalloc_base)
ffffca0000000000	-23 TB	ffffca7fffffffff	1 TB	... unused hole
ffffcb0000000000	-22 TB	ffffcb7fffffffff	1 TB	virtual memory map (vmemmap_base)
ffffcc0000000000	-21 TB	ffffcc7fffffffff	1 TB	... unused hole
ffffcd0000000000	-20 TB	ffffcd7fffffffff	16 TB	KASAN shadow memory
ffffce0000000000	-4 TB	ffffce7fffffffff	2 TB	... unused hole
ffffcf0000000000	-2 TB	ffffcf7fffffffff	2 TB	vaddr_end for KASLR
ffffd00000000000	-2 TB	ffffd7fffffffff	0.5 TB	cpu_entry_area mapping
ffffd80000000000	-1.5 TB	ffffd87fffffffff	0.5 TB	LDT remap for PTI
ffffd90000000000	-1 TB	ffffd97fffffffff	0.5 TB	%esp fixup stacks
				Identical layout to the 47-bit one from here on:
fffffe0000000000	-512 GB	fffffe07fffffffff	444 GB	... unused hole
fffffe0800000000	-68 GB	fffffe087fffffffff	64 GB	EFI region mapping space
fffffe0900000000	-4 GB	fffffe097fffffffff	2 GB	... unused hole
fffffe0a00000000	-2 GB	fffffe0a7fffffffff	512 MB	kernel text mapping, mapped to physical address 0
fffffe0b00000000	-2048 MB	fffffe0b7fffffffff	512 MB	kernel text mapping, mapped to physical address 0
fffffe0c00000000	-1536 MB	fffffe0c7fffffffff	1520 MB	module mapping space
fffffe0d00000000	-16 MB	fffffe0d7fffffffff	1520 MB	module mapping space
FIXADDR_START	~-11 MB	fffffe0d75fffff	~0.5 MB	kernel-internal fixmap range, variable size and offset
fffffe0e00000000	-10 MB	fffffe0e6007fff	4 kB	legacy vsyscall ABI
fffffe0e00000000	-2 MB	fffffe0e6007fff	2 MB	... unused hole

https://www.kernel.org/doc/Documentation/x86/x86_64/mm.txt

Execution Order

- Animated Slides (Power point)

Meltdown and Spectre

- Animated Slides (Power point)

Summary

- Vulnerabilities and Exploits
- Meltdown and Spectre
- Memory model and process Virtual Memory
- Side channeling
- Execution order – out-of-order execution
- Speculative execution

Acknowledgment

- IA-32 Intel: Architecture Software Developer's Manual Volume 3: System Programming Guide (Document 253668): Chapter 3 4.
- Meltdown Spectre for normal people (<https://github.com/neuhalje>)
- White paper on Meltdown Spectre (<http://cert-mu.govmu.org/English/Documents/White%20Papers/MELTDOWN%20-%20CERTMU%20WHITEPAPER.pdf>)