

Computer Architecture
Tutorial 2

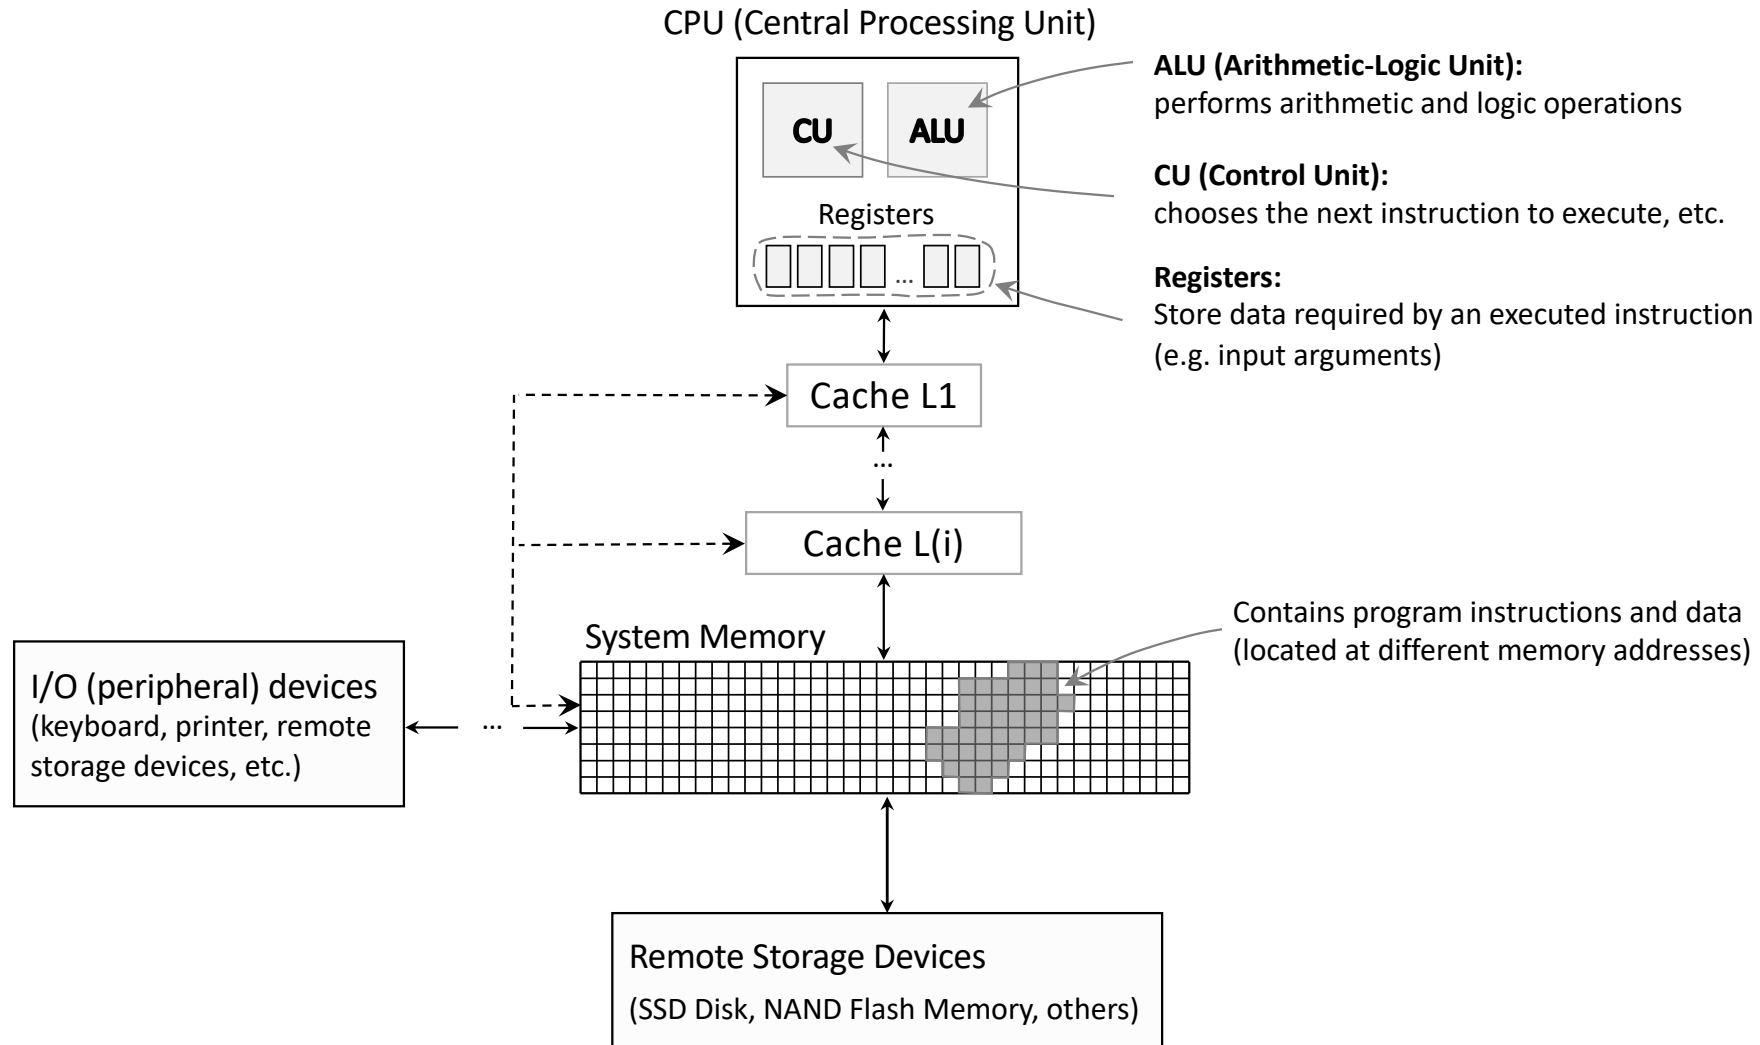
Introduction to Fundamental Ideas of Computer Architecture

Artem Burmyakov, Alexander Tormasov

September 02, 2021

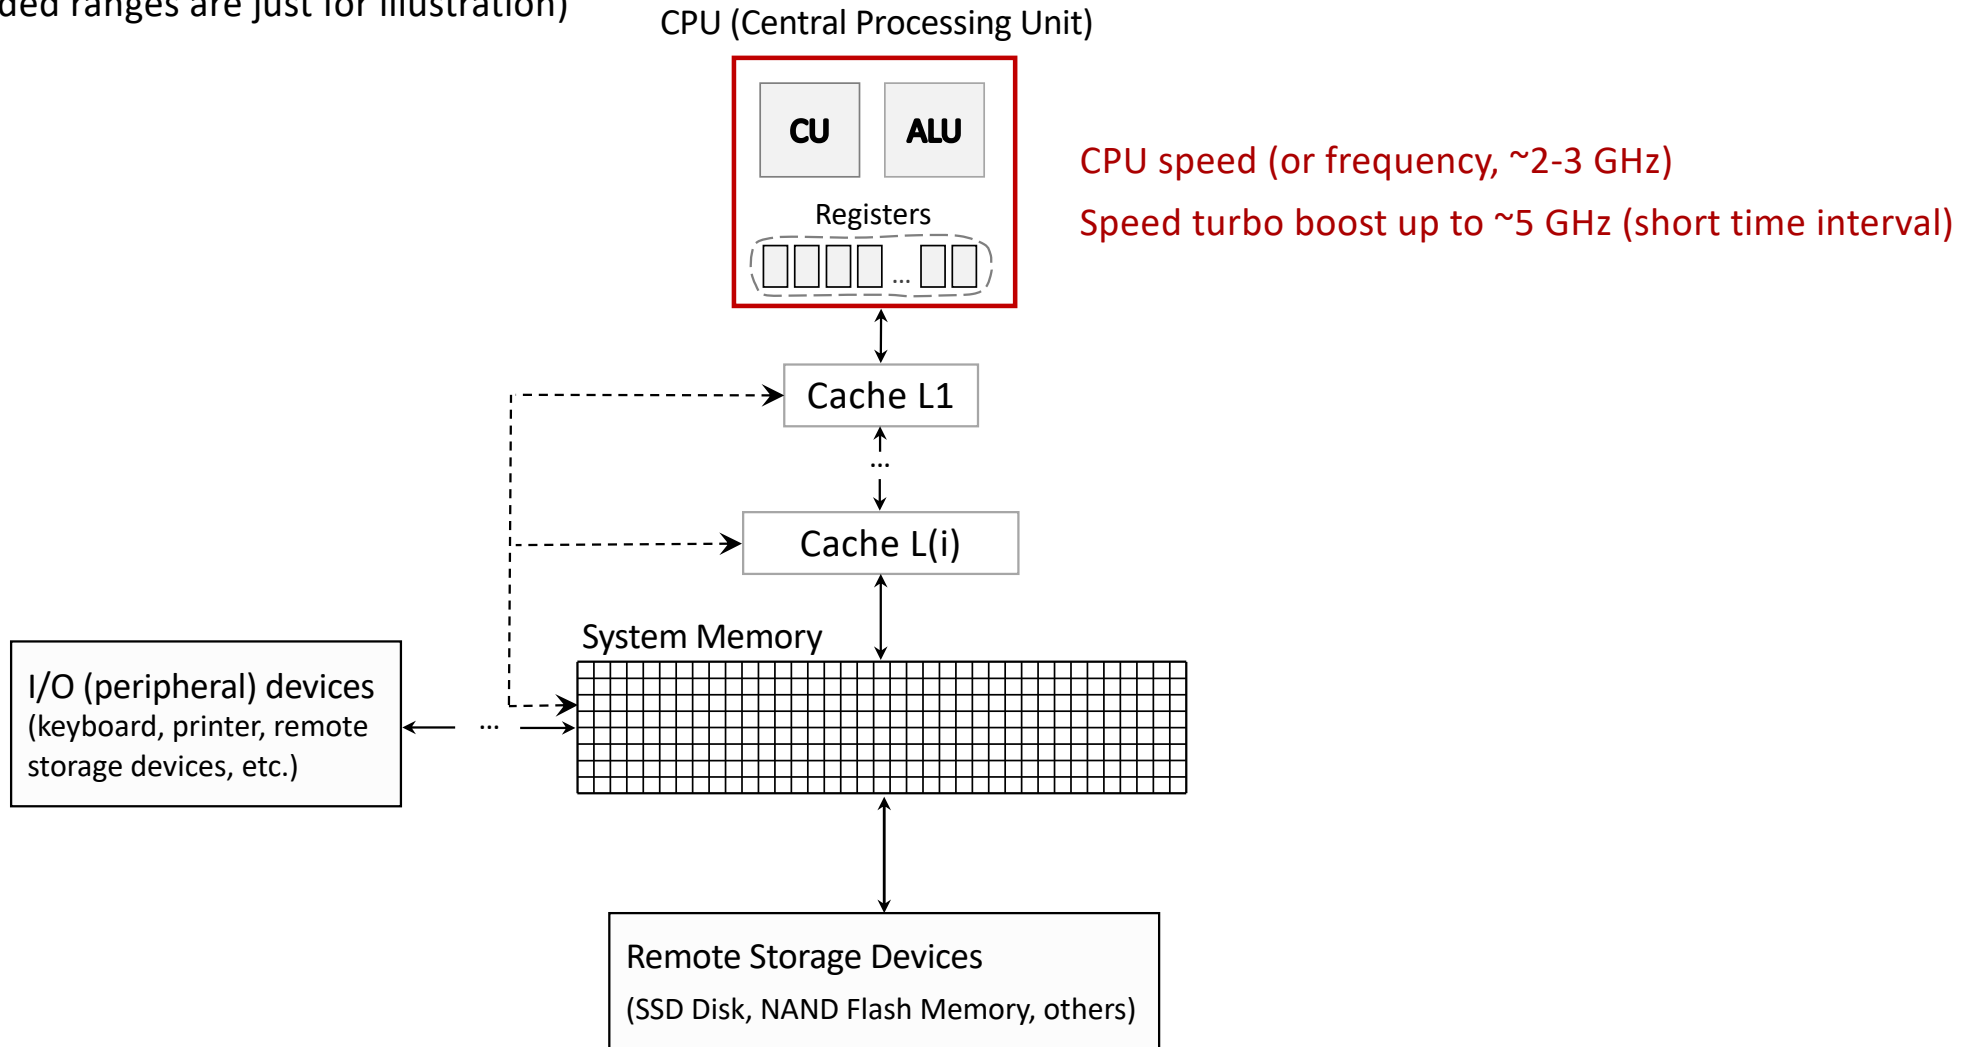


Recap of Key Computer Components



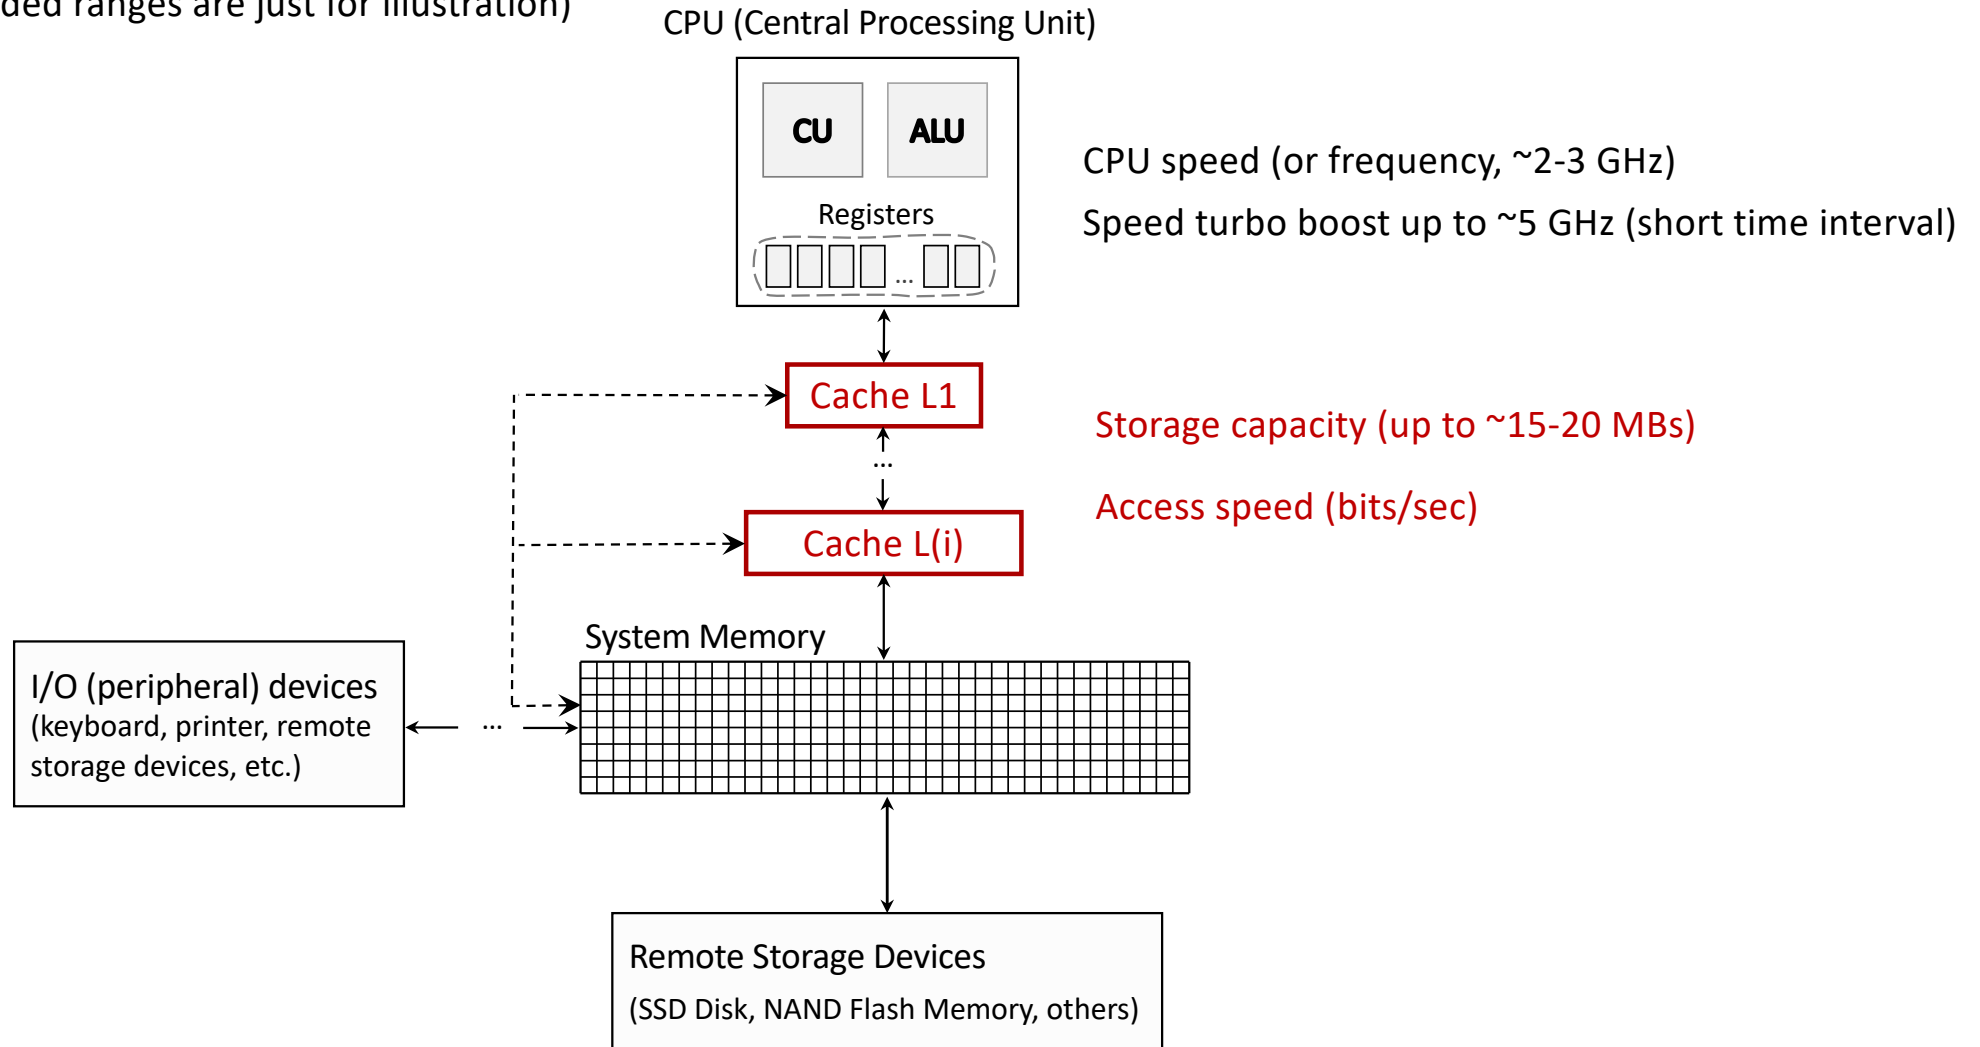
Some Performance Metrics

(provided ranges are just for illustration)



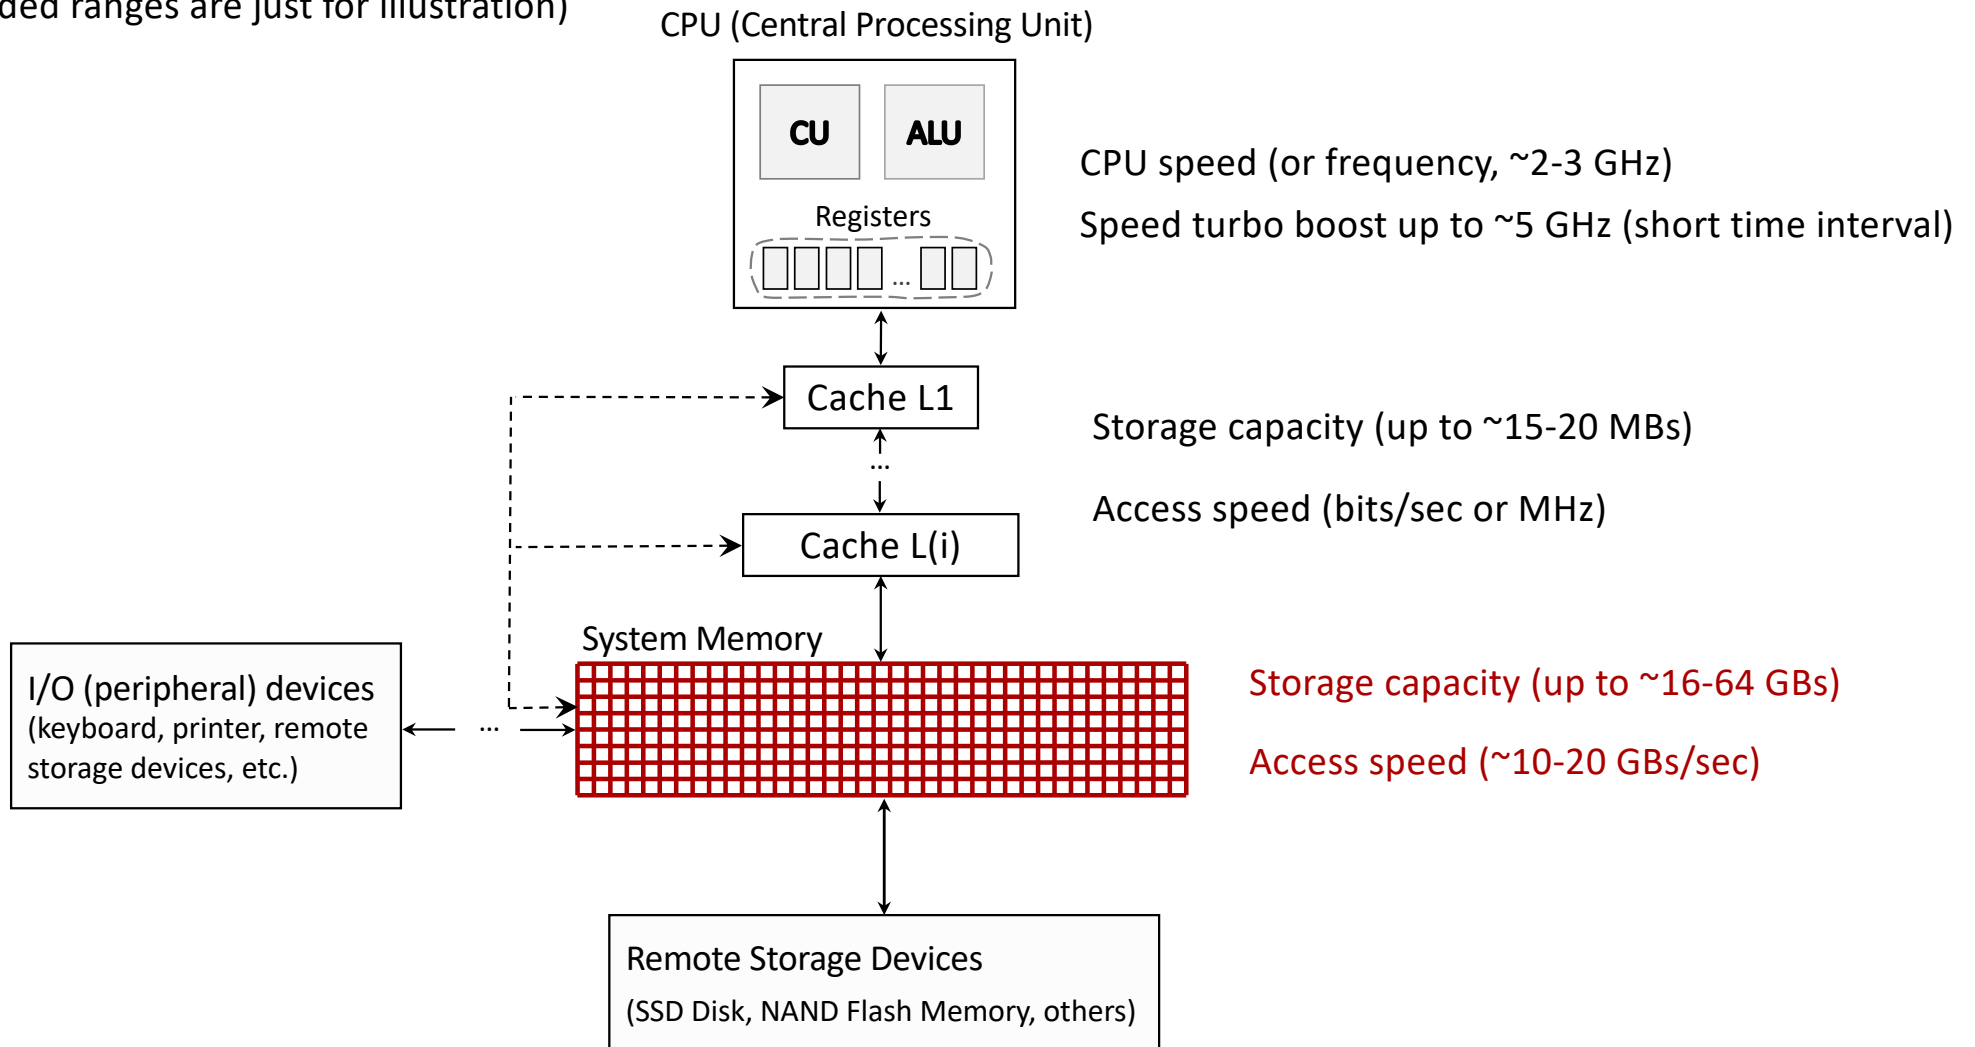
Some Performance Metrics

(provided ranges are just for illustration)



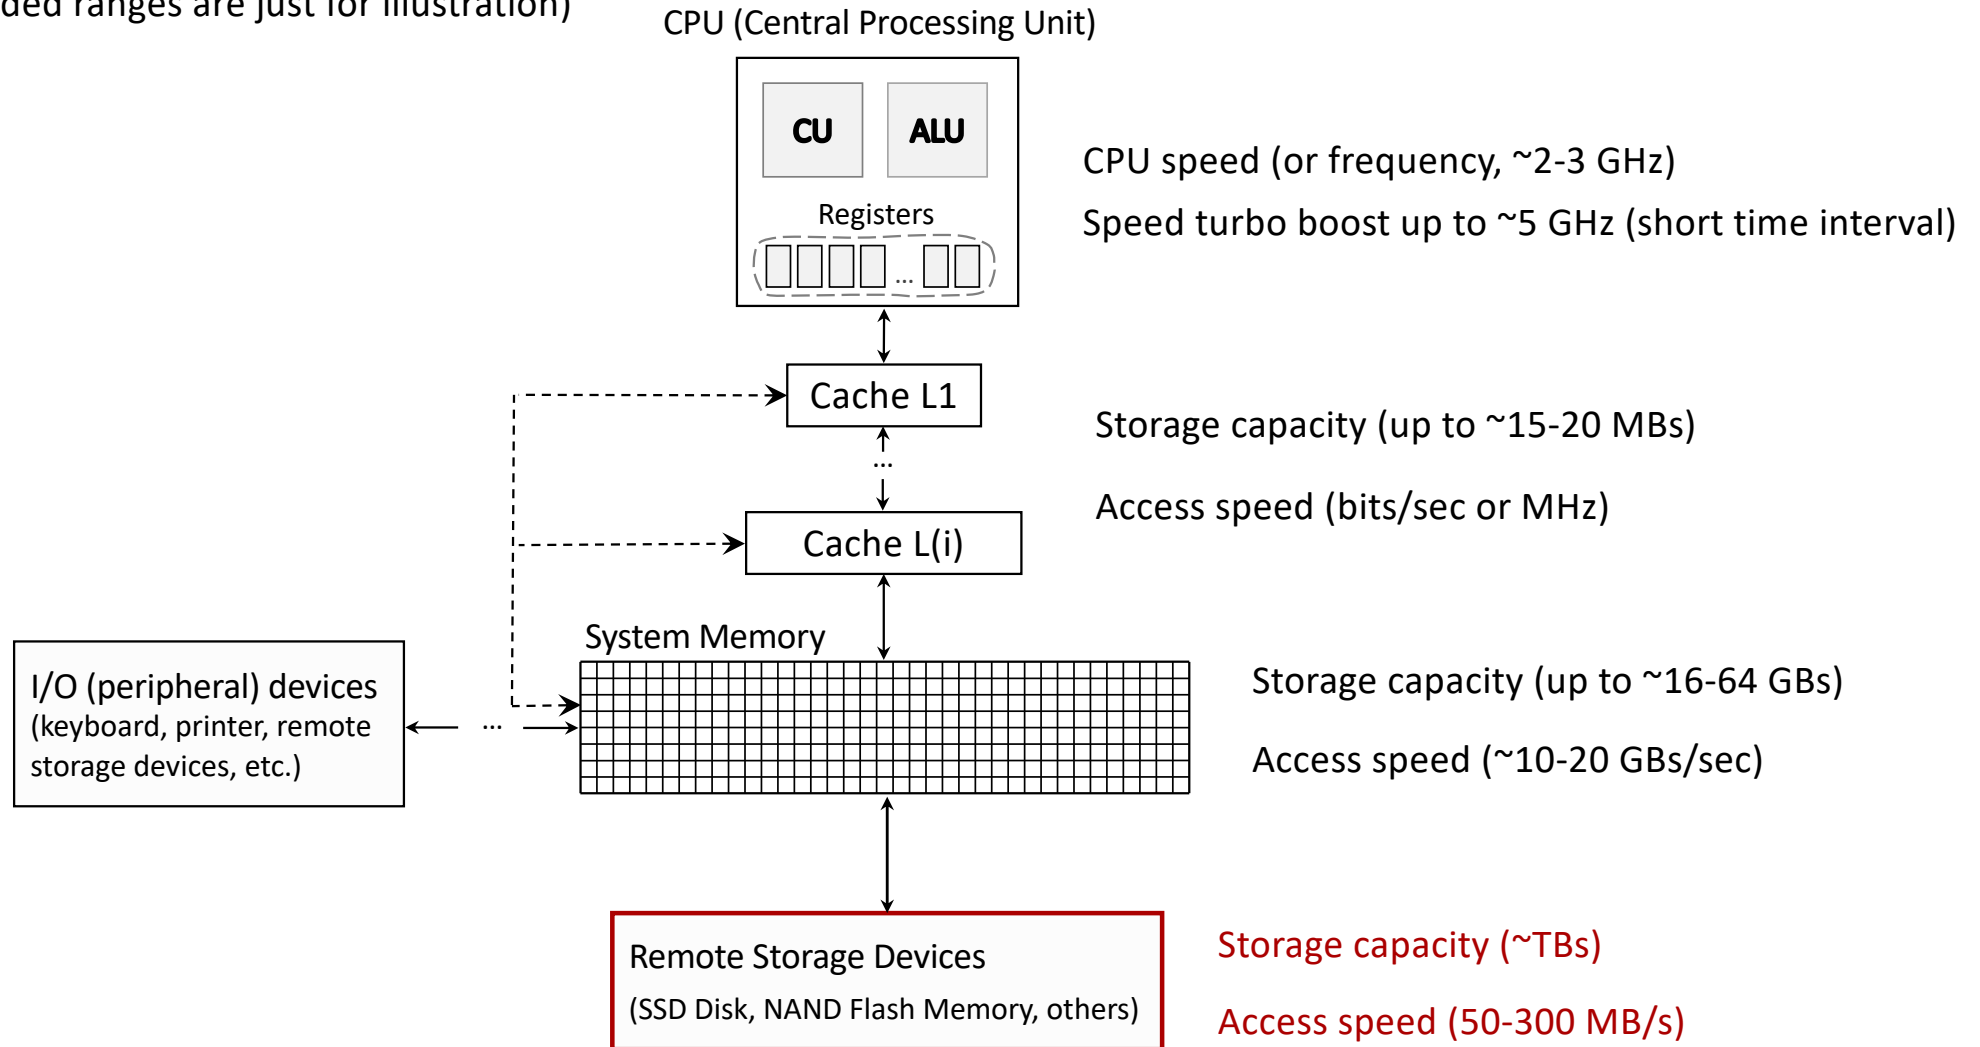
Some Performance Metrics

(provided ranges are just for illustration)



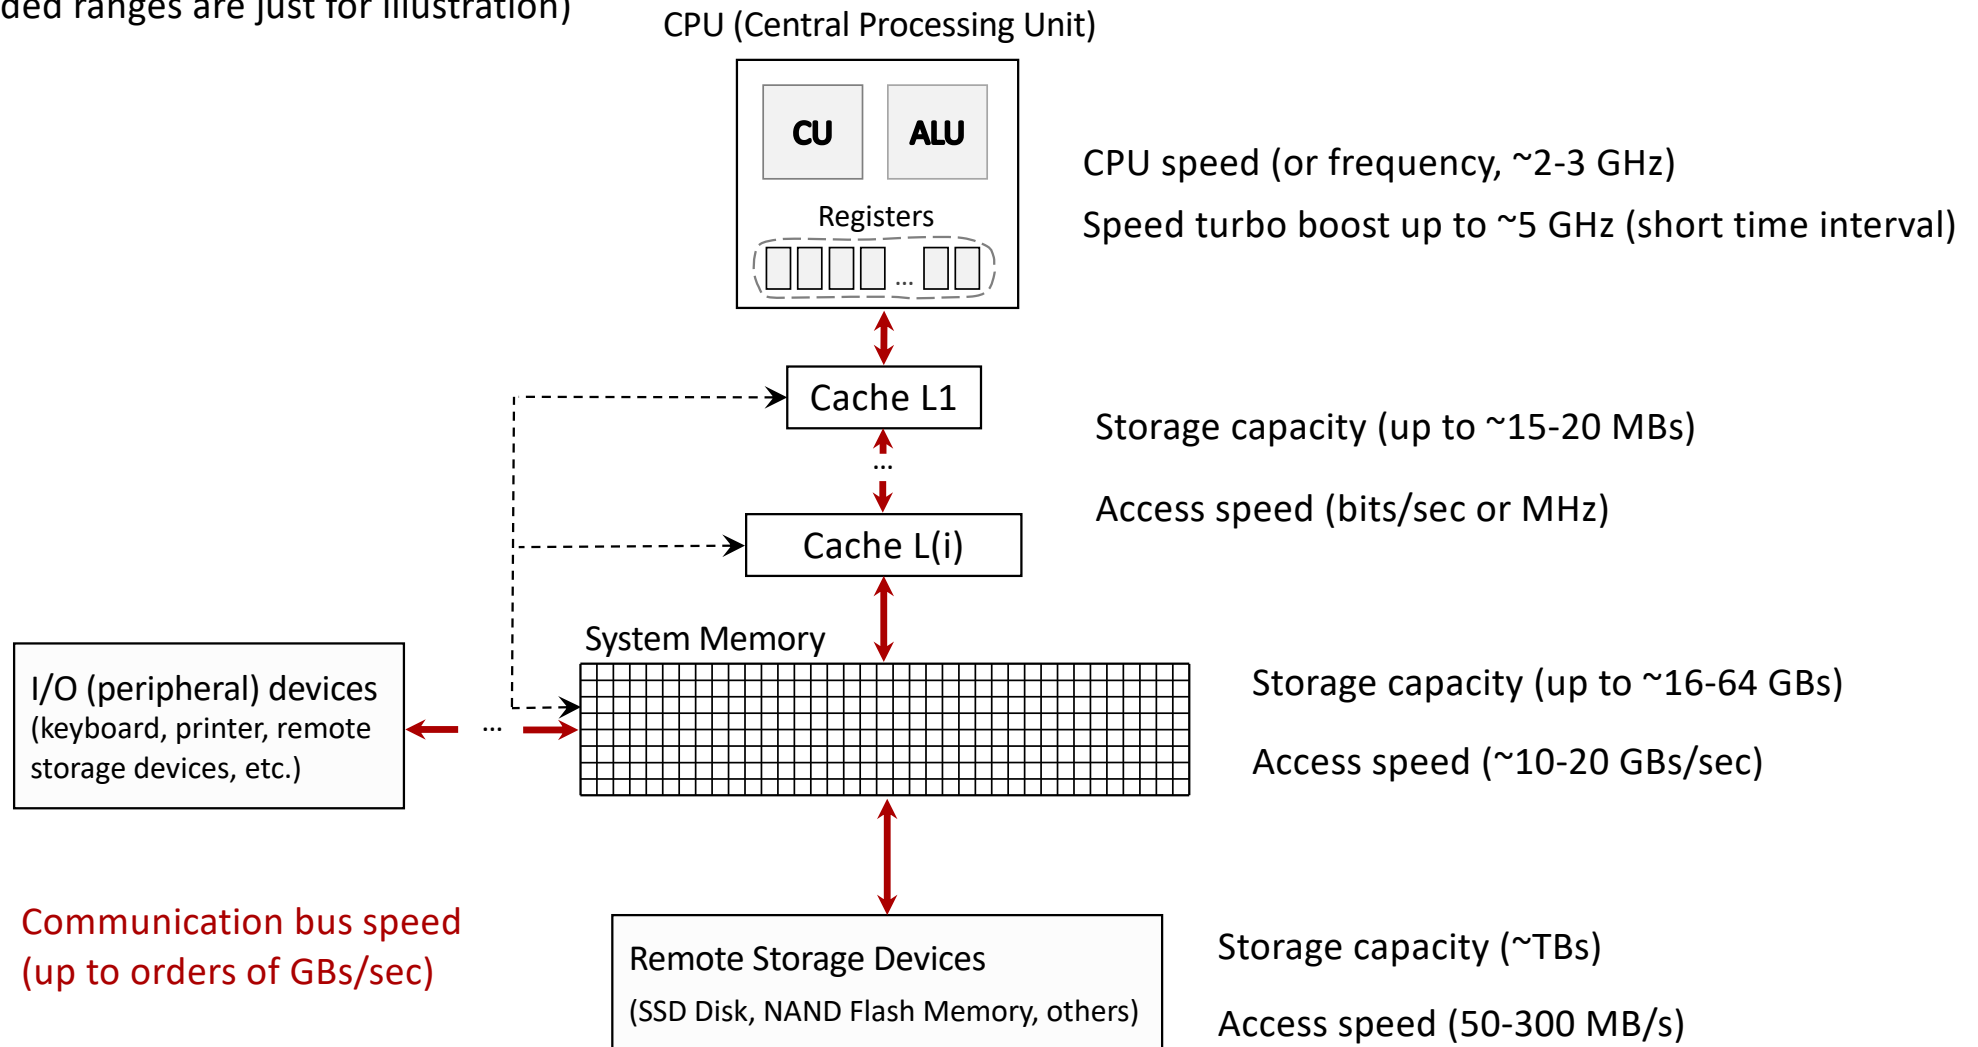
Some Performance Metrics

(provided ranges are just for illustration)



Some Performance Metrics

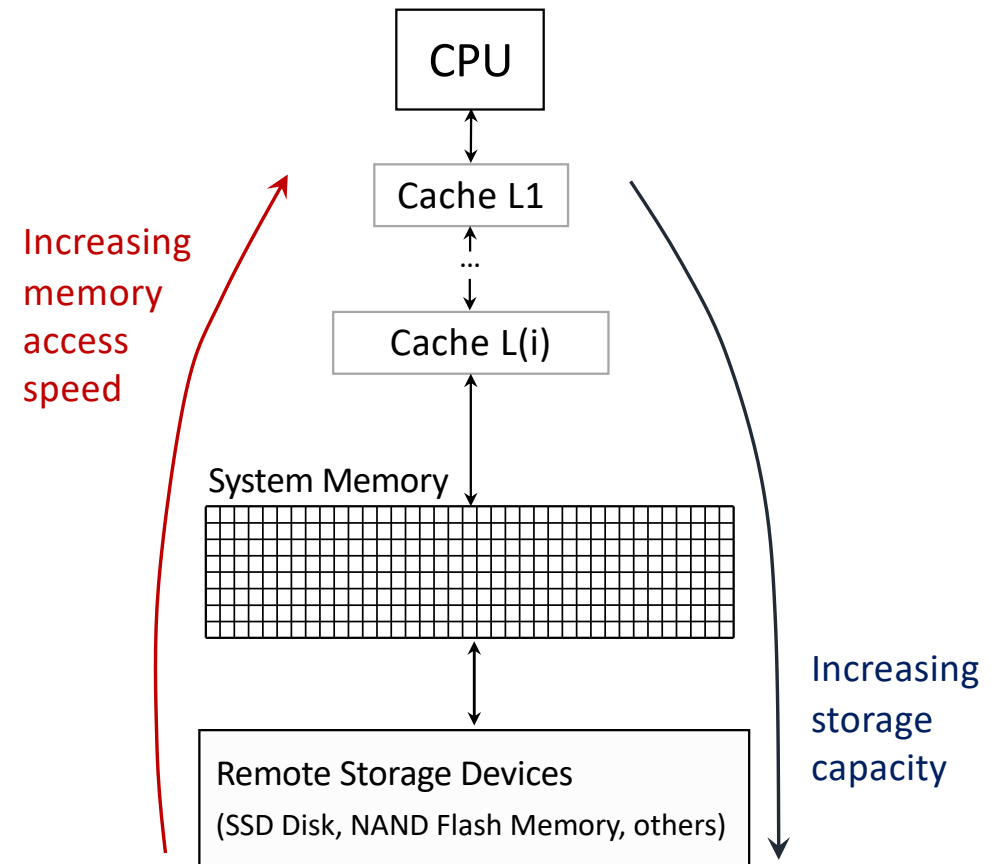
(provided ranges are just for illustration)



Fundamental (“Great”) Ideas of Computer Architecture

Fundamental (“Great”) Ideas of Computer Architecture

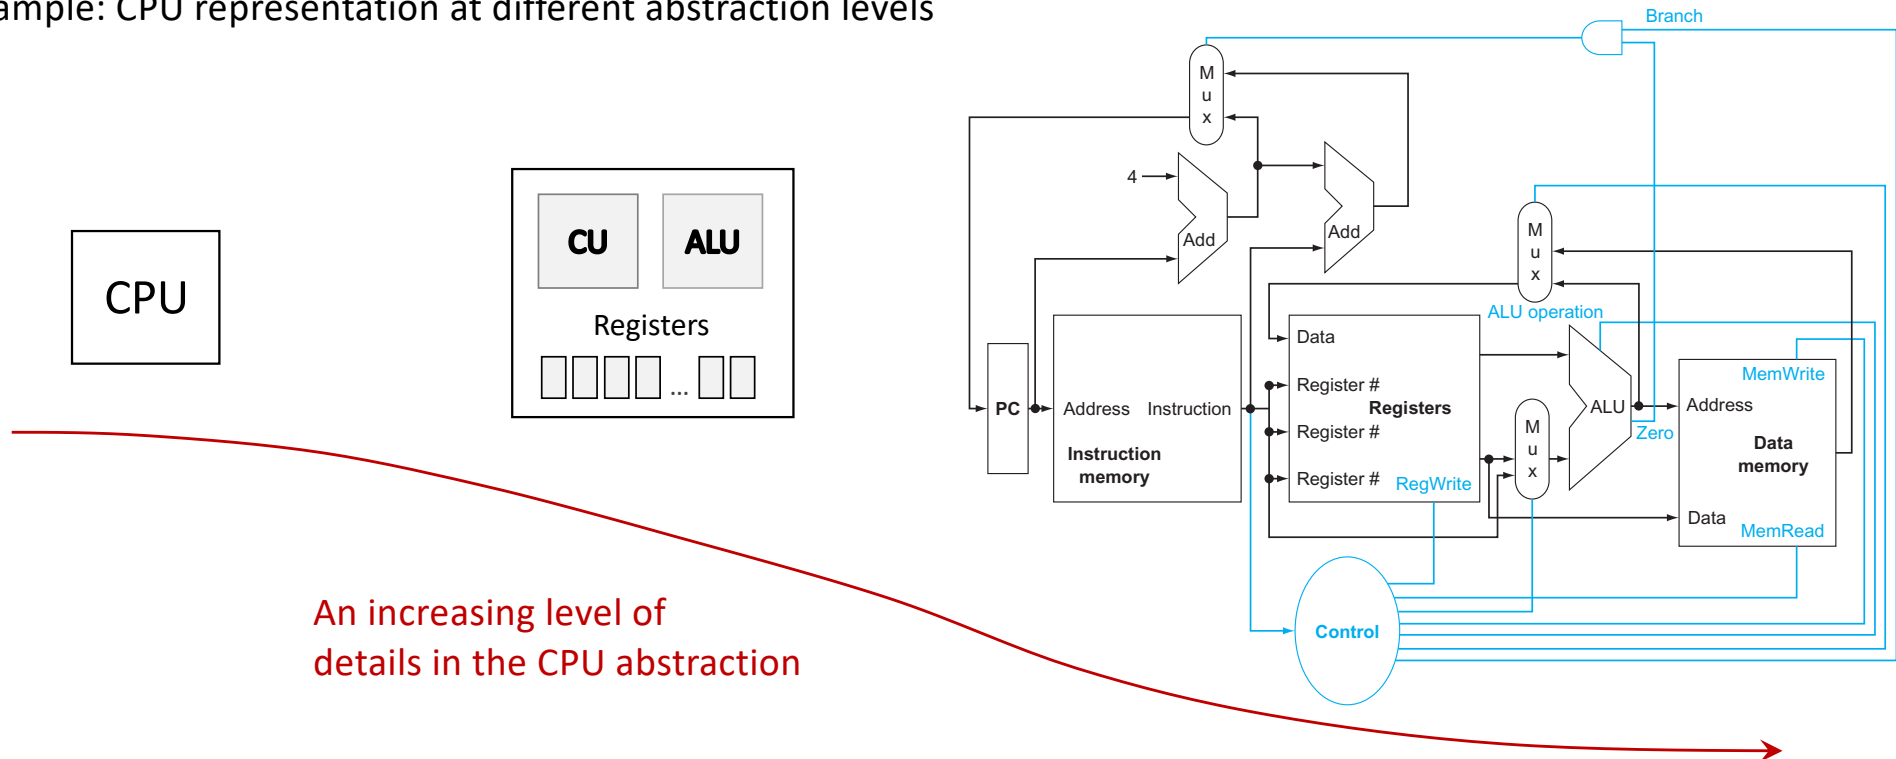
1. Hierarchy of memories



Fundamental (“Great”) Ideas of Computer Architecture

1. Hierarchy of memories
2. Use abstraction to simplify design

Example: CPU representation at different abstraction levels



Fundamental (“Great”) Ideas of Computer Architecture

1. Hierarchy of memories
2. Use abstraction to simplify design
3. Design for Moore's law


The number of transistors on a CPU chip doubles every 18-24 months (and thus, the CPU speed increases);

This observation is useful for long term projects:
By the time you complete your project, the CPU
speed might increase significantly

Note:
There is an ongoing discussion,
if Moore's law still holds

Moore's Law – The number of transistors on integrated circuit chips (1971-2018)

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important as other aspects of technological progress – such as processing speed or the price of electronic products – are linked to Moore's law.




Data source: Wikipedia (https://en.wikipedia.org/wiki/Transistor_count)

The data visualization is available at [OurWorldinData.org](https://ourworldindata.org). There you find more visualizations and research on this topic.

Licensed under [CC-BY-SA](#) by the author Max Roser.

Moore's Law Stagnation


Microprocessors



CPU speed and single-thread performance seem not to increase since ~2008

Moore's Law Stagnation

Microprocessors

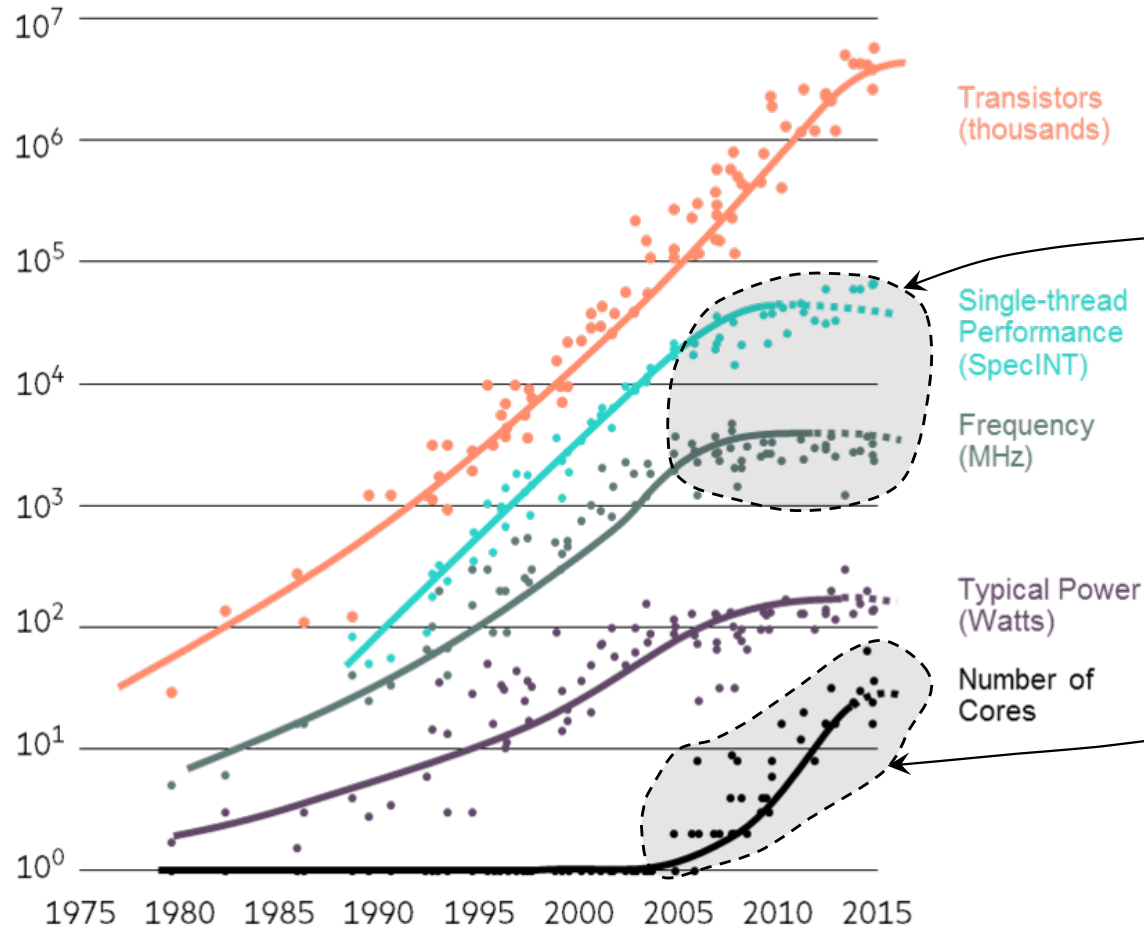


CPU speed and single-thread performance seem not to increase since ~2008

The number of cores increases instead

Moore's Law Stagnation

Microprocessors




CPU speed and single-thread performance seem not to increase since ~2008

The number of cores increases instead

Nowadays, the increase of performance is achieved through the number of CPUs (or CPU cores), rather than a further increase of CPU speed

Moore's Law Stagnation

Microprocessors




Heat dissipation problem –
the key reason for the CPU speed stagnation:

a higher clock rate →
a higher consumed power →
a higher power loss →
CPU overheating

Moore's Law Stagnation

Microprocessors



Heat dissipation problem –

the key reason for the CPU speed stagnation:

a higher clock rate →

a higher consumed power →

a higher power loss →


CPU overheating

The limitation of the speed of light – another problem:

signals in a CPU already travel nearly with a speed of light, between input and output pins;

Moore's Law Stagnation

Microprocessors



Heat dissipation problem –

the key reason for the CPU speed stagnation:

a higher clock rate →

a higher consumed power →

a higher power loss →

CPU overheating

The limitation of the speed of light – another problem:

signals in a CPU already travel nearly with a speed of light, between input and output pins;

There is a strong need for multiprocessor systems;

Multiprocessor systems however lead to various concurrency problems, such as race conditions

Fundamental (“Great”) Ideas of Computer Architecture

1. Hierarchy of memories
2. Use abstraction to simplify design
3. Design for Moore’s law
4. Performance via parallelism

Fundamental (“Great”) Ideas of Computer Architecture

4. Performance via parallelism


Computer program
(a set of instructions to be executed);
loaded into system memory

Instruction 1
Instruction 2
Instruction 3
Instruction 4
Instruction 5
Instruction 6
Instruction 7
...

Assumptions:

- Instructions are executed sequentially;
- CPU executes one instruction at a time (not always a case)

Single processor (or uniprocessor) system:



One instruction is
executed at a time
(no pipelining is assumed)

Fundamental (“Great”) Ideas of Computer Architecture

4. Performance via parallelism


Computer program

(a set of instructions to be executed);
loaded into system memory

Instruction 1
Instruction 2
Instruction 3
Instruction 4
Instruction 5
Instruction 6
Instruction 7
...

Step 1: Instruction
loaded into CPU

Single processor (or uniprocessor) system:



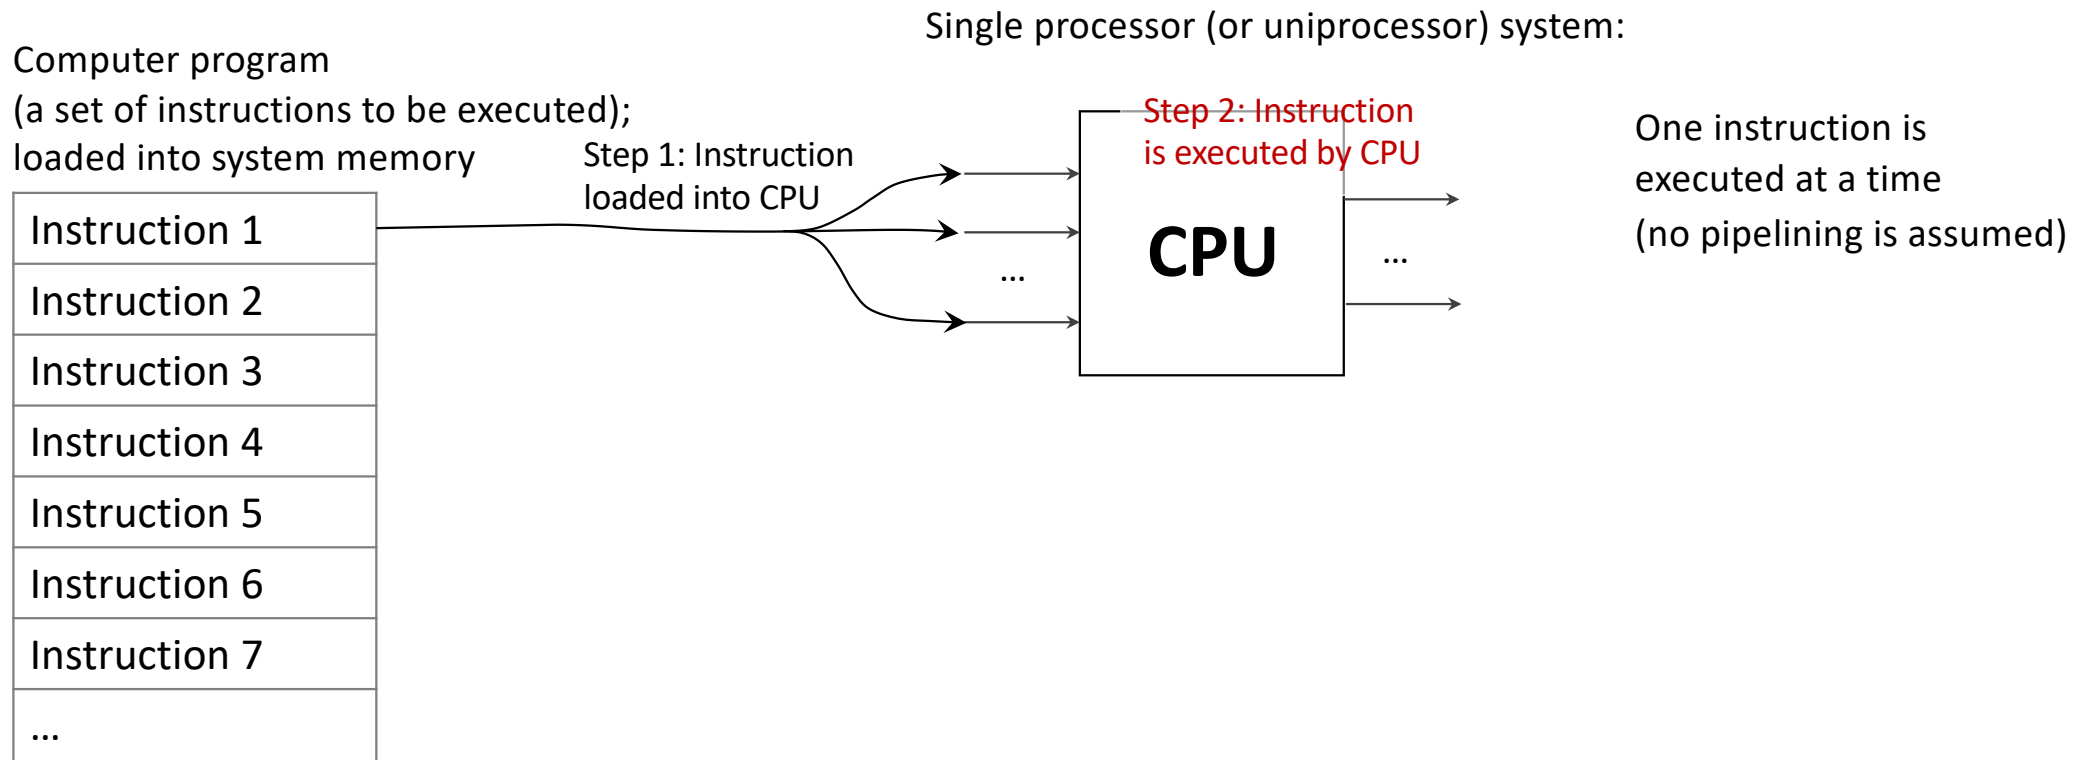
One instruction is
executed at a time
(no pipelining is assumed)

Assumptions:

- Instructions are executed sequentially;
- CPU executes one instruction at a time (not always a case)

Fundamental (“Great”) Ideas of Computer Architecture

4. Performance via parallelism



Assumptions:

- Instructions are executed sequentially;
- CPU executes one instruction at a time (not always a case)

Fundamental (“Great”) Ideas of Computer Architecture

4. Performance via parallelism

Computer program

(a set of instructions to be executed);
loaded into system memory

Instruction 1
Instruction 2
Instruction 3
Instruction 4
Instruction 5
Instruction 6
Instruction 7
...

Step 1: Instruction
loaded into CPU

Single processor (or uniprocessor) system:

Step 2: Instruction
is executed by CPU

CPU

One instruction is
executed at a time
(no pipelining is assumed)

Step 3: Result is
stored into memory

Assumptions:

- Instructions are executed sequentially;
- CPU executes one instruction at a time (not always a case)

Fundamental (“Great”) Ideas of Computer Architecture

4. Performance via parallelism

Computer program

(a set of instructions to be executed);
loaded into system memory

Instruction 1
Instruction 2
Instruction 3
Instruction 4
Instruction 5
Instruction 6
Instruction 7
...

Step 1: Instruction
loaded into CPU

Single processor (or uniprocessor) system:

Step 2: Instruction
is executed by CPU

CPU

Step 3: Result is
stored into memory

One instruction is
executed at a time
(no pipelining is assumed)

Each of these steps take some time!

Assumptions:

- Instructions are executed sequentially;
- CPU executes one instruction at a time (not always a case)

Fundamental (“Great”) Ideas of Computer Architecture

4. Performance via parallelism


Computer program
(a set of instructions to be executed);
loaded into system memory

Instruction 1
Instruction 2
Instruction 3
Instruction 4
Instruction 5
Instruction 6
Instruction 7
...

Assumptions:

- Instructions are executed sequentially;
- CPU executes one instruction at a time (not always a case)

Multiprocessor (or multicore) system:



M instructions can be
executed at a time,
simultaneously

Fundamental (“Great”) Ideas of Computer Architecture

4. Performance via parallelism


Computer program
(a set of instructions to be executed);
loaded into system memory

Instruction 1
Instruction 2
Instruction 3
Instruction 4
Instruction 5
Instruction 6
Instruction 7
...

Assumptions:

- Instructions are executed sequentially;
- CPU executes one instruction at a time (not always a case)

Multiprocessor (or multicore) system:



M instructions can be
executed at a time,
simultaneously

Fundamental (“Great”) Ideas of Computer Architecture


4. Performance via parallelism

Computer program

(a set of instructions to be executed);
loaded into system memory

Instruction 1
Instruction 2
Instruction 3
Instruction 4
Instruction 5
Instruction 6
Instruction 7
...

Multiprocessor (or multicore) system:



M instructions can be
executed at a time,
simultaneously

Assumptions:

- Instructions are executed sequentially;
- CPU executes one instruction at a time (not always a case)

Fundamental (“Great”) Ideas of Computer Architecture


4. Performance via parallelism

Computer program
(a set of instructions to be executed);
loaded into system memory

Instruction 1
Instruction 2
Instruction 3
Instruction 4
Instruction 5
Instruction 6
Instruction 7
...

Simultaneous
execution
(restrictions apply)

Multiprocessor (or multicore) system:



M instructions can be
executed at a time,
simultaneously

Multiprocessing aims at
speeding-up program
execution, or simultaneous
execution of multiple
programs

Assumptions:


- Instructions are executed sequentially;
- CPU executes one instruction at a time (not always a case)

Fundamental (“Great”) Ideas of Computer Architecture

4. Performance via parallelism


Computer program

(a set of instructions to be executed);
loaded into system memory



Simultaneous
execution
(restrictions apply)

Multiprocessor (or multicore) system:



Assumptions:

- Instructions are executed sequentially;
- CPU executes one instruction at a time (not always a case)

M instructions can be
executed at a time,
simultaneously

Multiprocessing aims at
speeding-up program
execution, or simultaneous
execution of multiple
programs

The major problem:
Some instructions depend on
the execution result of
previous instructions, and,
thus, have to wait for their
completion

Fundamental (“Great”) Ideas of Computer Architecture

4. Performance via parallelism


Computer program

(a set of instructions to be executed);
loaded into system memory

Instruction 1
Instruction 2
Instruction 3
Instruction 4
Instruction 5
Instruction 6
Instruction 7
...

Simultaneous
execution
(restrictions apply)

Multiprocessor (or multicore) system:



M instructions can be
executed at a time,
simultaneously

Multiprocessing aims at
speeding-up program
execution, or simultaneous
execution of multiple
programs

**Many problems remain open;
These problems got a lot of
attention after Moore's law
“slow down”**

Assumptions:

- Instructions are executed sequentially;
- CPU executes one instruction at a time (not always a case)


Fundamental (“Great”) Ideas of Computer Architecture

1. Hierarchy of memories
2. Use abstraction to simplify design
3. Design for Moore’s law
4. Performance via parallelism
5. Performance via pipelining

Fundamental (“Great”) Ideas of Computer Architecture

5. Performance via pipelining

Our custom CPU



Instruction Set, supported by our CPU:

Operation	Operation Code
Logical “OR”	0
Logical “AND”	1

Instruction Format:

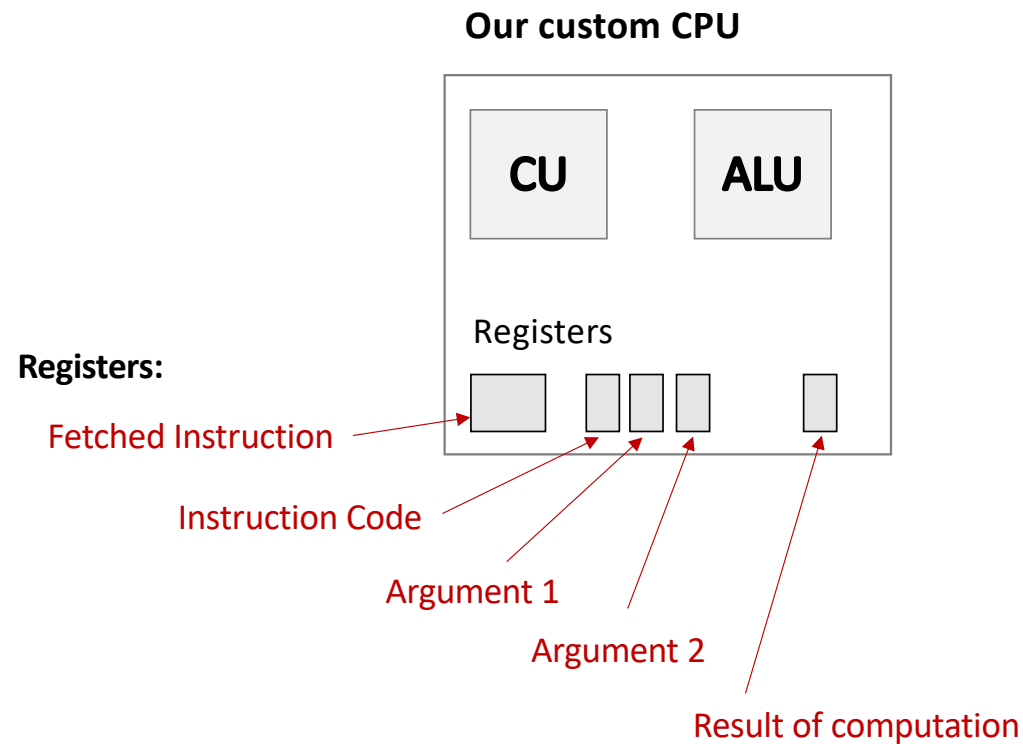
Operation Code	Argument 1	Argument 2
----------------	------------	------------

Sample Instruction:

0	1	0
---	---	---

Fundamental (“Great”) Ideas of Computer Architecture

5. Performance via pipelining



Instruction Set, supported by our CPU:

Operation	Operation Code
Logical “OR”	0
Logical “AND”	1

Instruction Format:

Operation Code	Argument 1	Argument 2
----------------	------------	------------

Sample Instruction:

0	1	0
---	---	---


Fundamental (“Great”) Ideas of Computer Architecture

5. Performance via pipelining

Instructions in System Memory

Instruction 1: 010
Instruction 2
Instruction 3
Instruction 4
Instruction 5
...

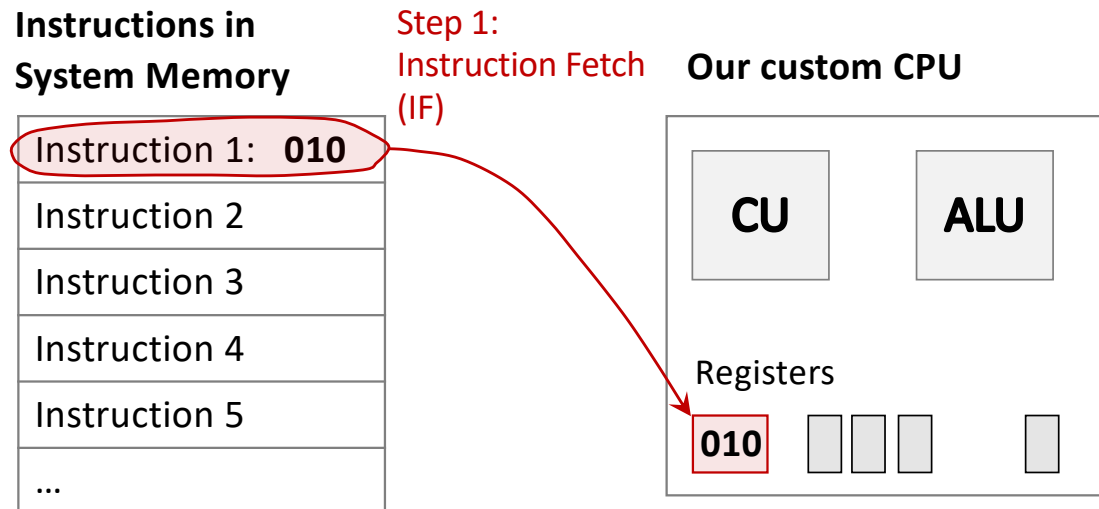
Our custom CPU



Fundamental (“Great”) Ideas of Computer Architecture

5. Performance via pipelining

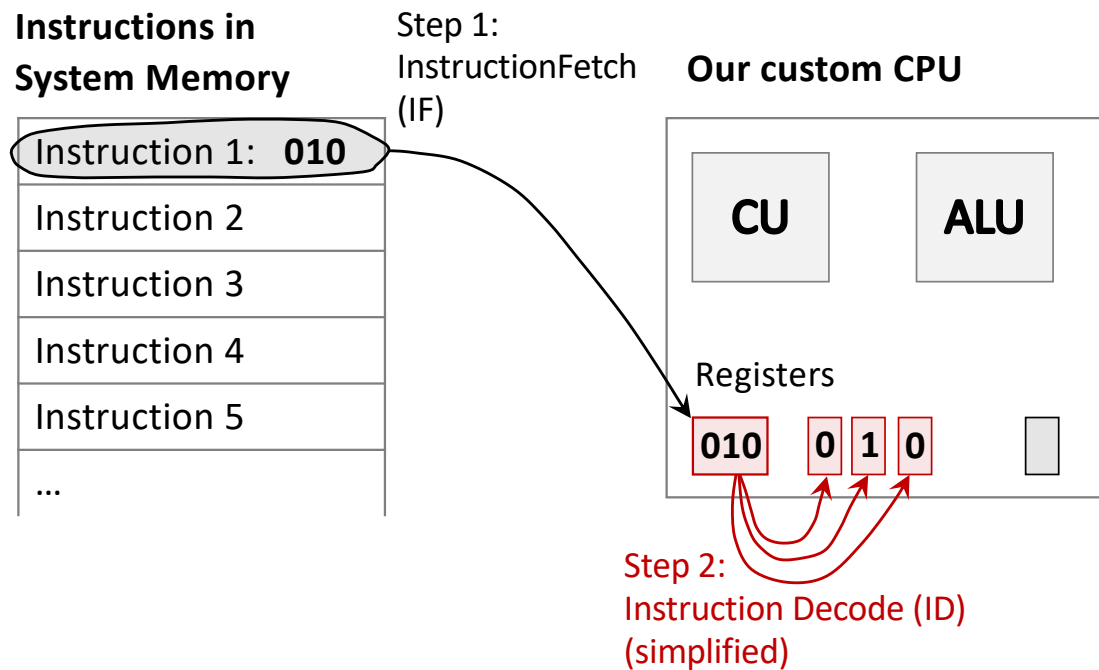
Each instruction is executed in several steps:



Fundamental (“Great”) Ideas of Computer Architecture

5. Performance via pipelining

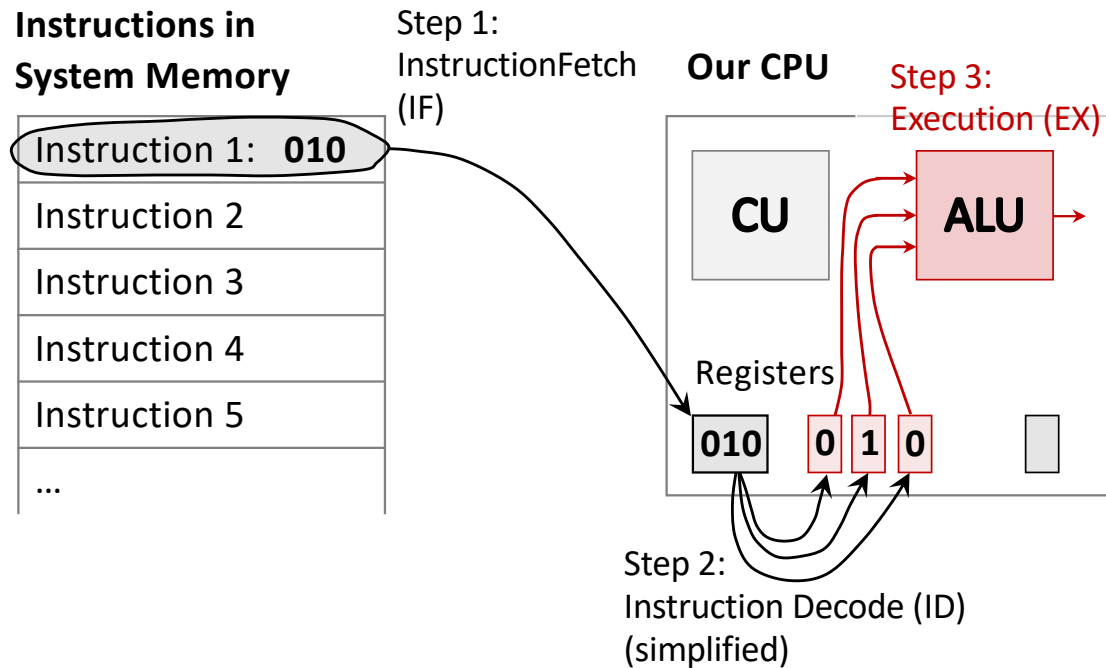
Each instruction is executed in several steps:



Fundamental (“Great”) Ideas of Computer Architecture

5. Performance via pipelining

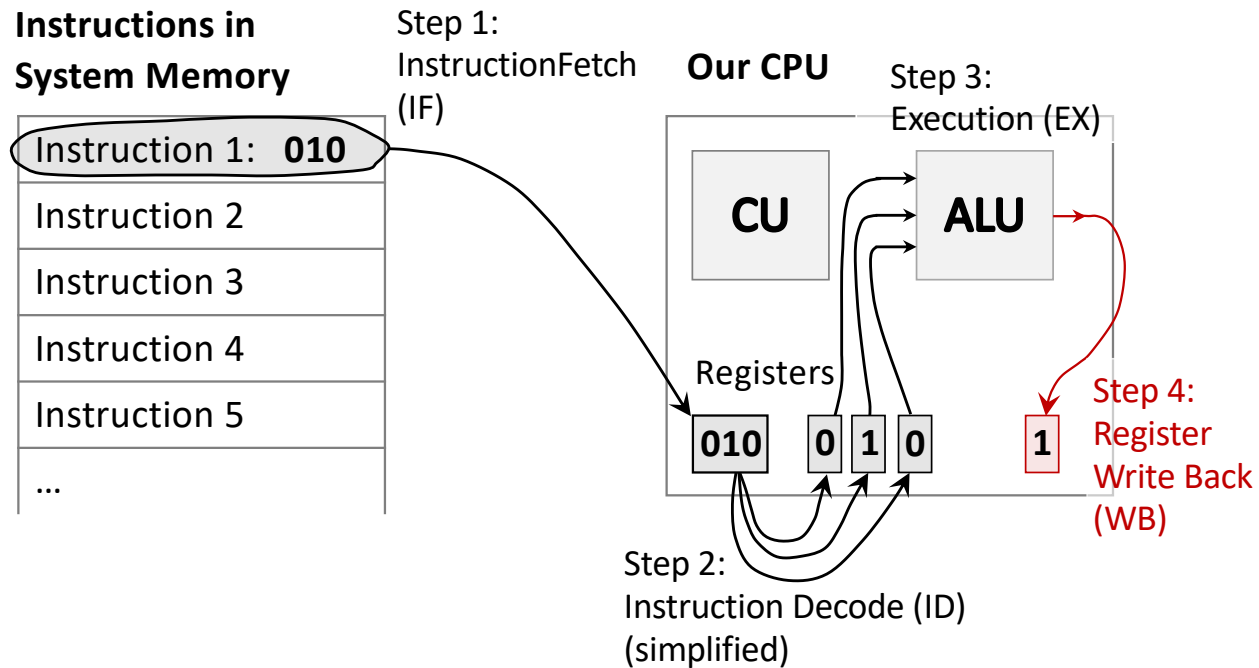
Each instruction is executed in several steps:



Fundamental (“Great”) Ideas of Computer Architecture

5. Performance via pipelining

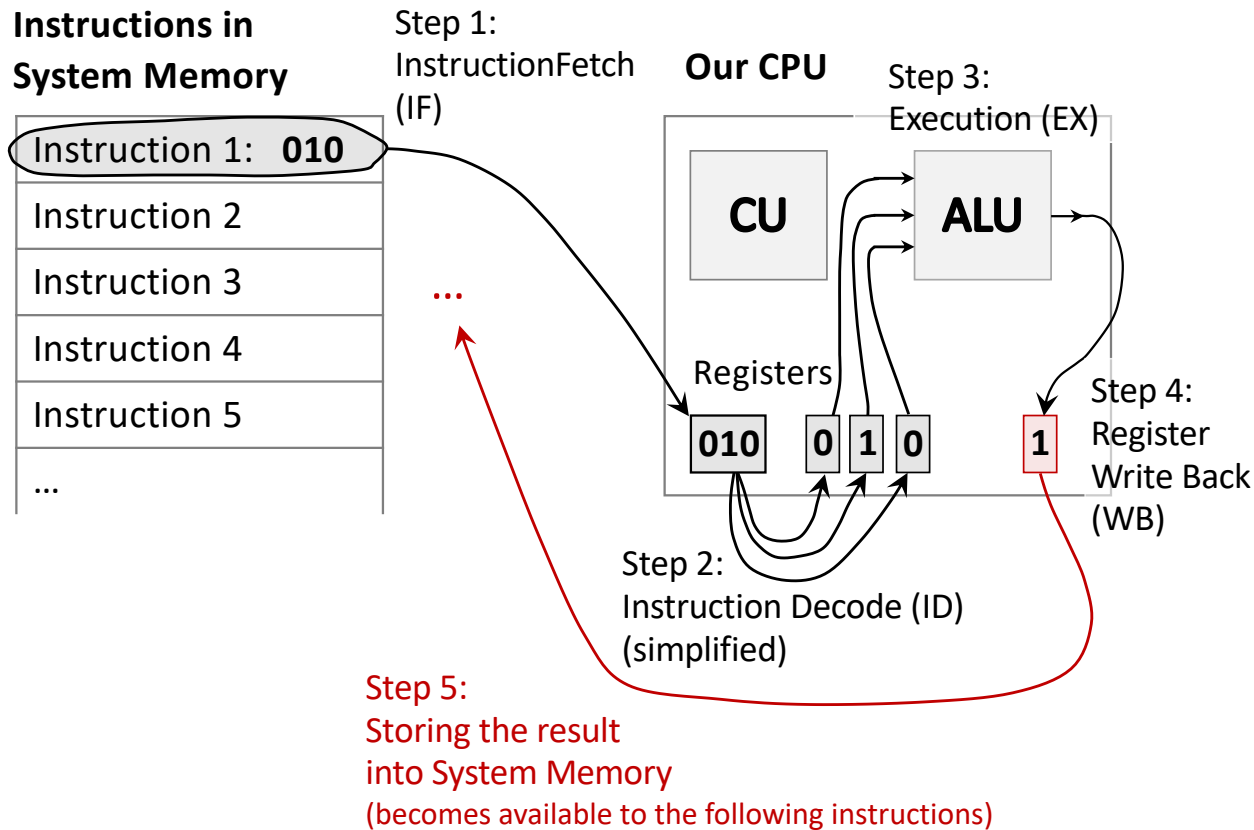
Each instruction is executed in several steps:



Fundamental (“Great”) Ideas of Computer Architecture

5. Performance via pipelining

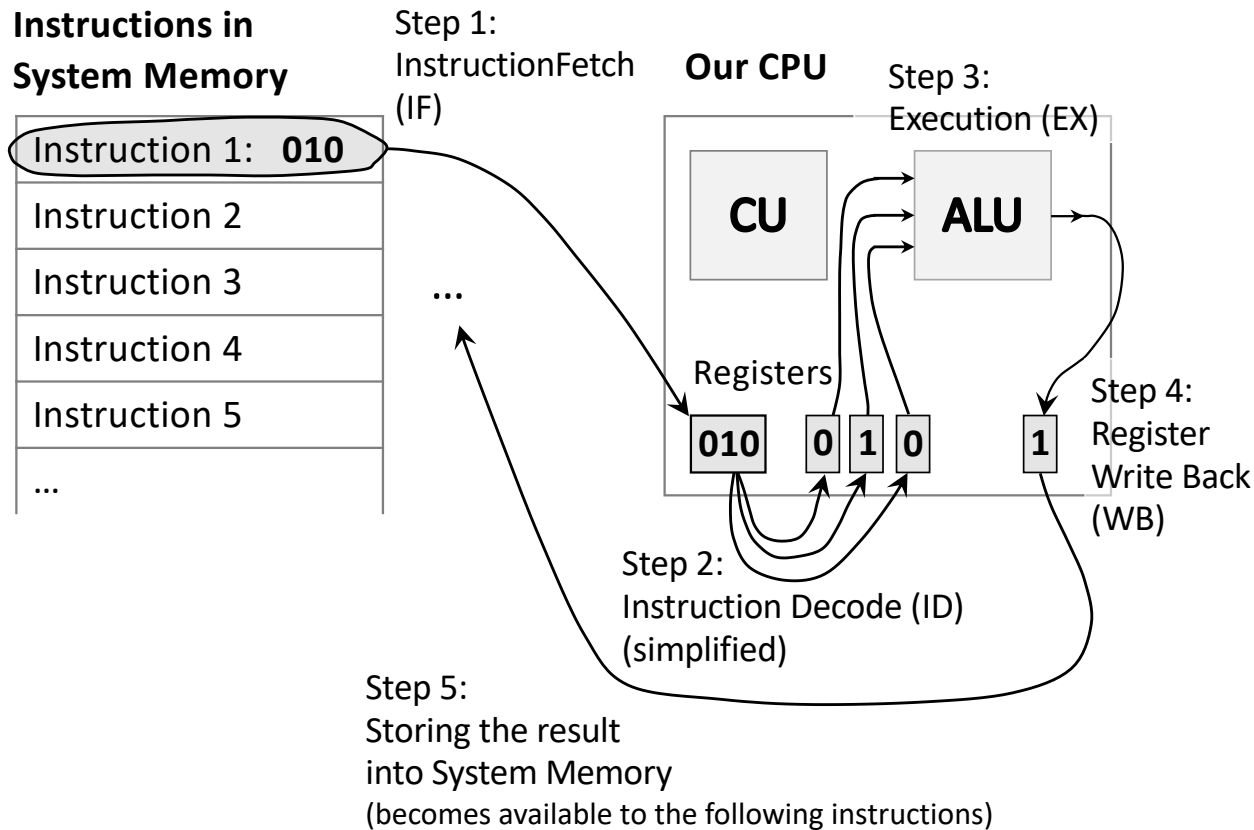
Each instruction is executed in several steps:



Fundamental (“Great”) Ideas of Computer Architecture

5. Performance via pipelining

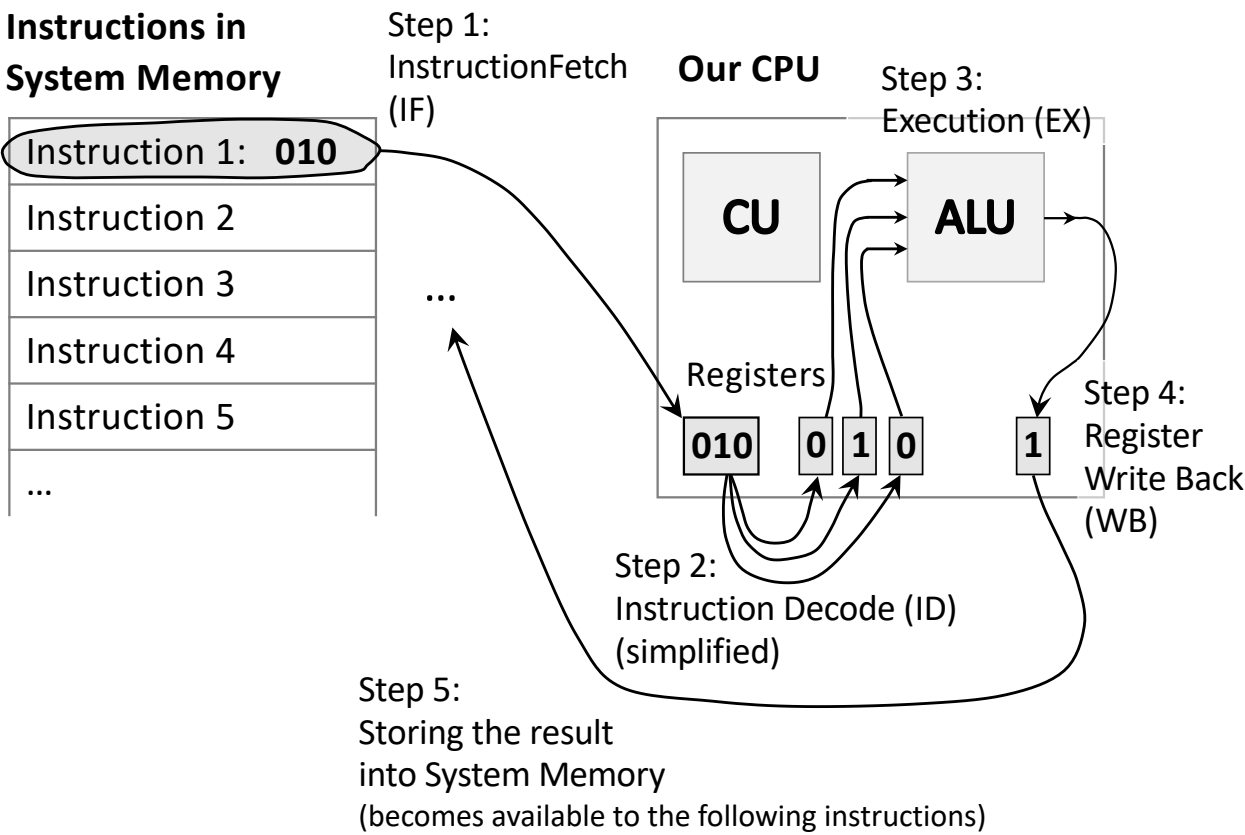
Each instruction is executed in several steps:



Fundamental (“Great”) Ideas of Computer Architecture

5. Performance via pipelining

Each instruction is executed in several steps:



Idea of pipelining:

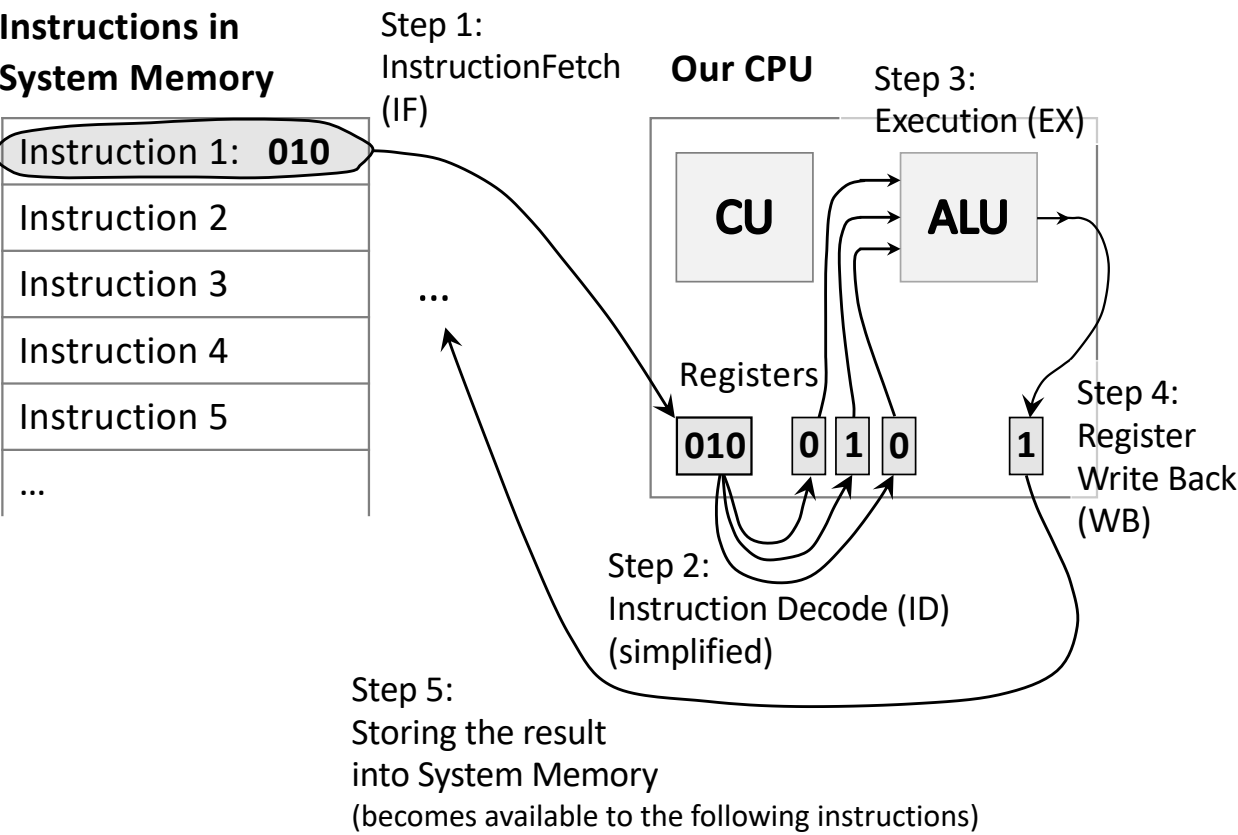
CPU might execute several instructions simultaneously, but each instruction at a different execution stage

Instr. No. \ Clock cycle	1	2	3	4	5	6	7
1							
2							
3							
4							
5							

Fundamental (“Great”) Ideas of Computer Architecture

5. Performance via pipelining

Each instruction is executed in several steps:



Idea of pipelining:

CPU might execute several instructions simultaneously, but each instruction at a different execution stage


Instr. No. \ Clock cycle	1	2	3	4	5	6	7
1	IF						
2							
3							
4							
5							

IF: Instruction Fetch

Fundamental (“Great”) Ideas of Computer Architecture

5. Performance via pipelining

Each instruction is executed in several steps:



Idea of pipelining:

CPU might execute several instructions simultaneously, but each instruction at a different execution stage

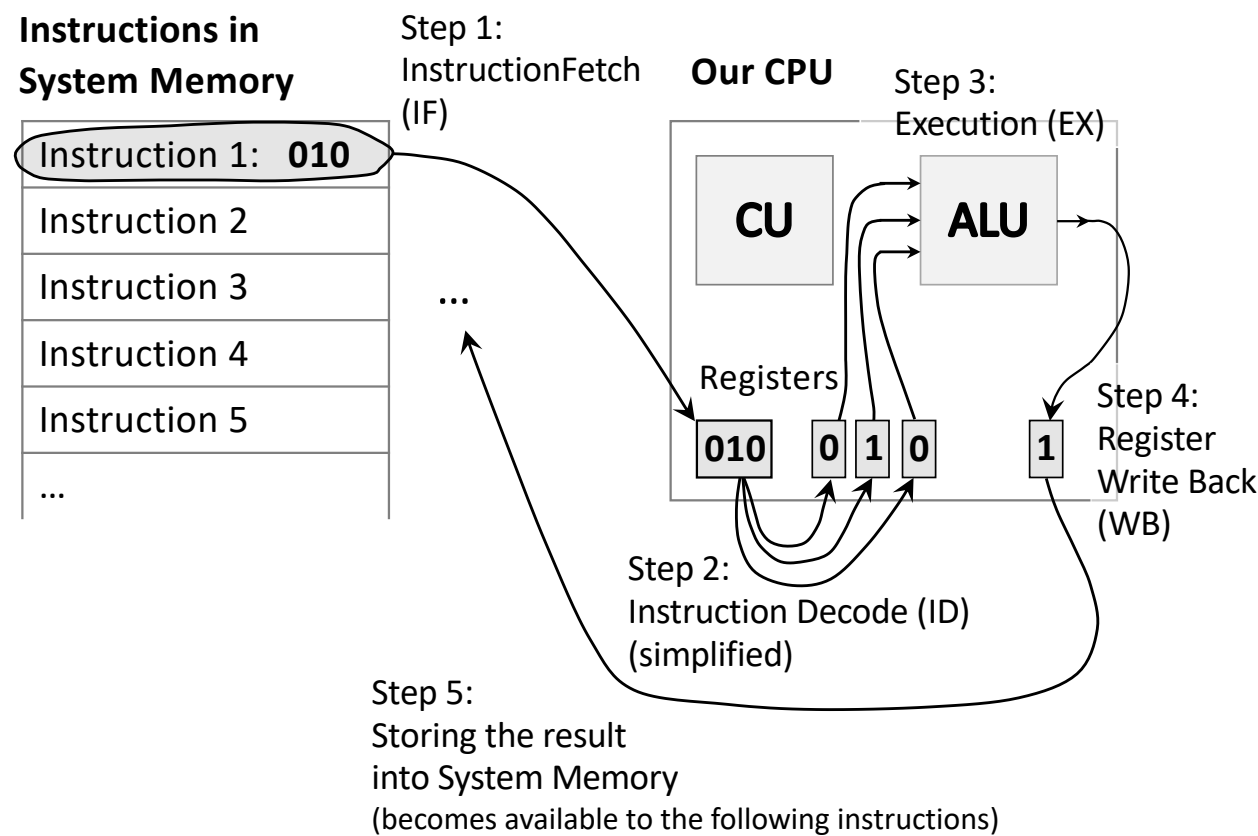
Instr. No. \ Clock cycle	1	2	3	4	5	6	7
1	IF	ID					
2		IF					
3							
4							
5							

IF: Instruction Fetch
ID: Instruction Decode

Fundamental (“Great”) Ideas of Computer Architecture

5. Performance via pipelining

Each instruction is executed in several steps:



Idea of pipelining:

CPU might execute several instructions simultaneously, but each instruction at a different execution stage


Instr. No. \ Clock cycle	1	2	3	4	5	6	7
1	IF	ID					
2		IF					
3							
4							
5							

IF: Instruction Fetch
ID: Instruction Decode

Fundamental (“Great”) Ideas of Computer Architecture

5. Performance via pipelining

Each instruction is executed in several steps:



Idea of pipelining:

CPU might execute several instructions simultaneously, but each instruction at a different execution stage

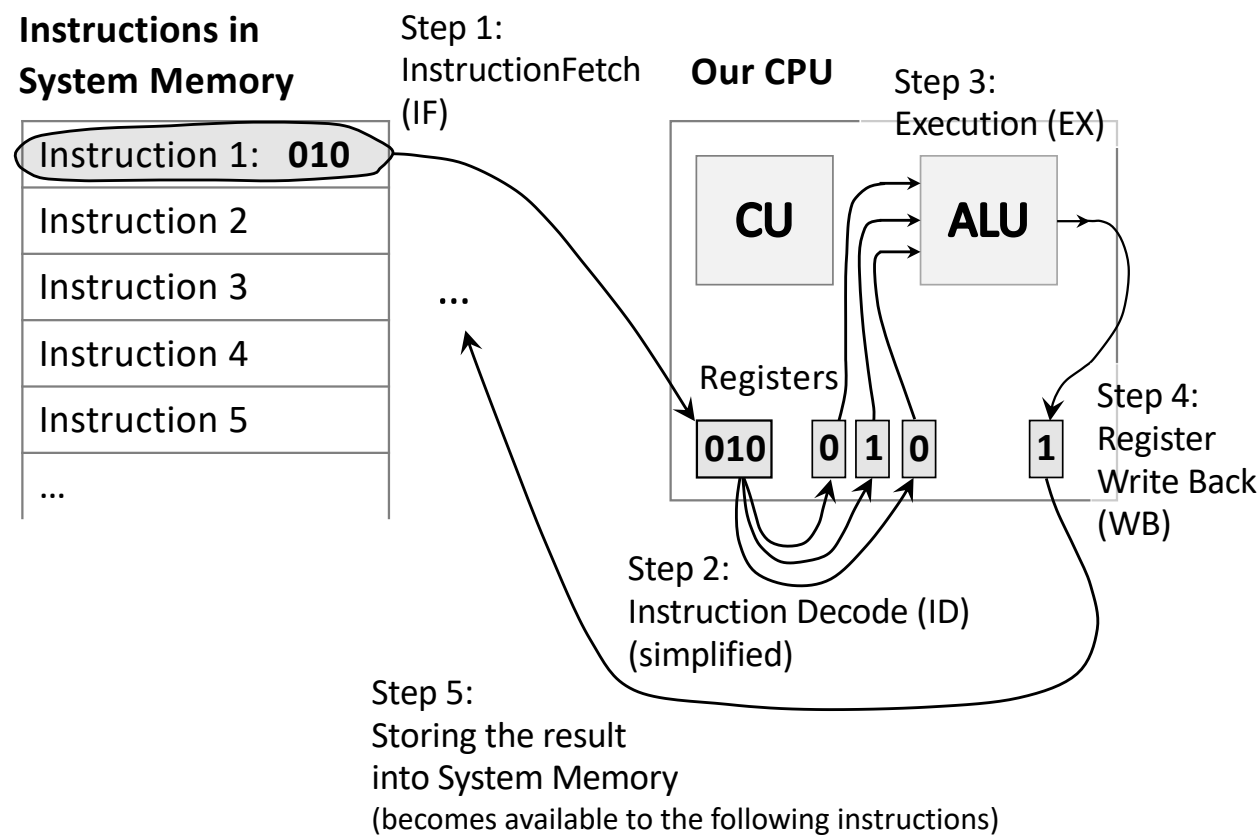
Instr. No. \ Clock cycle	1	2	3	4	5	6	7
1	IF	ID	EX				
2		IF	ID				
3			IF				
4							
5							

IF: Instruction Fetch
ID: Instruction Decode
EX: Execution

Fundamental (“Great”) Ideas of Computer Architecture

5. Performance via pipelining

Each instruction is executed in several steps:



Idea of pipelining:

CPU might execute several instructions simultaneously, but each instruction at a different execution stage

Instr. No. \ Clock cycle	1	2	3	4	5	6	7
1	IF	ID	EX	MEM	WB		
2		IF	ID	EX	MEM	WB	
3			IF	ID	EX	MEM	WB
4				IF	ID	EX	MEM
5					IF	ID	EX

IF: Instruction Fetch
ID: Instruction Decode
EX: Execution
MEM: Memory access
WB: register Write Back

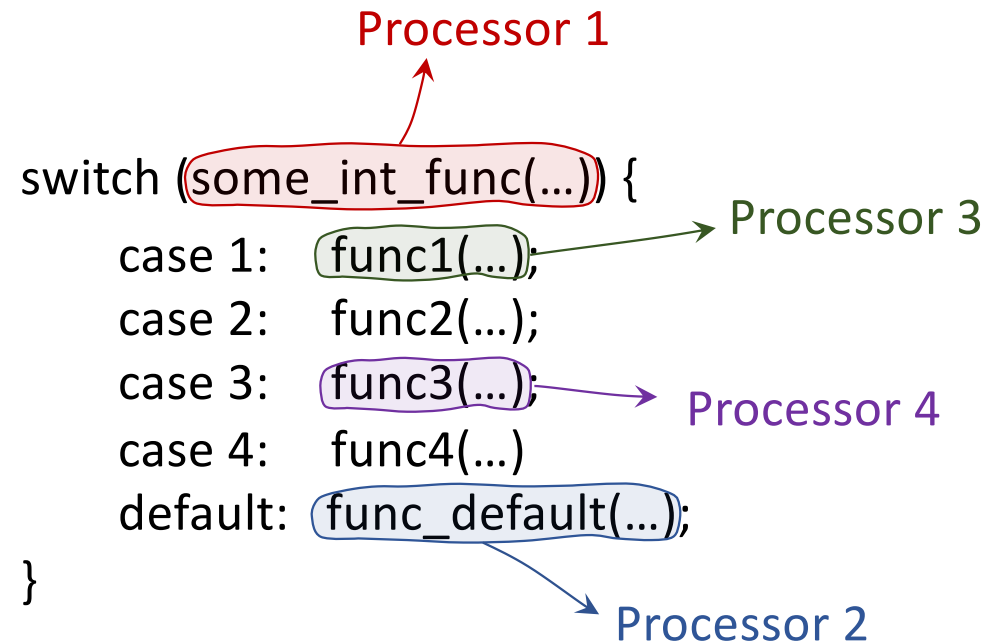
Fundamental (“Great”) Ideas of Computer Architecture

1. Hierarchy of memories
2. Use abstraction to simplify design
3. Design for Moore’s law
4. Performance via parallelism
5. Performance via pipelining
6. Performance via speculation (prediction)

C++ code sample example:

Assumptions:

- 4 processors are available;
- No dependencies between calling functions



For this specific code, we try to predict (“speculate”) 3 most possible cases to be executed

Fundamental (“Great”) Ideas of Computer Architecture

1. Hierarchy of memories
2. Use abstraction to simplify design
3. Design for Moore’s law
4. Performance via parallelism
5. Performance via pipelining
6. Performance via speculation (prediction)
7. Dependability (reliability) via redundancy

Some redundant (or “spare”) components are introduced (e.g. CPU or memory unit), to increase the reliability of a computer platform (e.g. in a spacecraft);

Fundamental (“Great”) Ideas of Computer Architecture

1. Hierarchy of memories
2. Use abstraction to simplify design
3. Design for Moore’s law
4. Performance via parallelism
5. Performance via pipelining
6. Performance via speculation (prediction)
7. Dependability (reliability) via redundancy
8. Make the common case fast

The computer platform should be optimized for the most common use case expected

Fundamental (“Great”) Ideas of Computer Architecture

1. Hierarchy of memories
2. Use abstraction to simplify design
3. Design for Moore’s law
4. Performance via parallelism
5. Performance via pipelining
6. Performance via speculation (prediction)
7. Dependability (reliability) via redundancy
8. Make the common case fast
9. State machines

State machine – a set of states and transitions between them; a convenient way to model systems behaviour

A sample state machine




Image is taken from <https://dwarves.foundation/n/finite-state-machine/>

Fundamental (“Great”) Ideas of Computer Architecture

1. Hierarchy of memories
2. Use abstraction to simplify design
3. Design for Moore’s law
4. Performance via parallelism
5. Performance via pipelining
6. Performance via speculation (prediction)
7. Dependability (reliability) via redundancy
8. Make the common case fast
9. State machines